

A Market World Model for Public LLMs: Complete, Honest, Auditable Crypto Cognition before Decision

Anonymous Author(s)

Abstract

Most quant stacks start from *strategy*. Public LLMs (GPT-, DeepSeek-, GLM-, and Gemini-class agents) that act in cryptocurrency markets fail earlier: they lack a complete, honest, auditable *market cognition base*. We present an anonymized *financial world-model runtime* that compiles asynchronous multi-band evidence into a point-in-time (PIT) world bundle. Epistemic observations $O_{j,t} = (x, \tau, q, g, r)$ and compilation Π_t yield $\mathcal{F}_t^{AI}$; completeness, honesty (B, U, H), and WMI/ACWMI support a typed abstention contract; evidence IDs make actions auditable. Four live public LLMs (temperature 0) on the *same* 100 OOS asset-days per arm show a three-arm behavior ladder: Compiled thin-world abstain 1.0; Ungated (same content, no hard flag) mean 0.68 (range 0.43–0.86); Raw 0.04–0.75. Disclosure helps but is vendor-dependent; the typed boolean enforces full cross-vendor refusal, and models that act under Ungated lose money on thin support. A non-trading grounding workflow shows the same asymmetry for analysis: with the compiled three-band bundle, models recover ready/missing sets and tilt signs near-perfectly (mean F1/acc. 0.97); from a raw feed the same scores collapse to 0.04–0.23. Economically, the compiled world state outperforms traditional baselines out-of-sample: the world-model content rule earns Sharpe 0.76 / CE +0.13 versus momentum 0.10 / –0.21 and buy-and-hold –1.40 / –1.16 (pre-specified $\Delta CE=0.334$, $p=0.034$), and live LLMs consuming the bundle avoid the losses that raw-feed agents incur (mean $\Delta CE +1.07$). We detail the runtime architecture (collectors, vintage store, BandPIT clock, bundle schema, LLM adapter) as a systems contribution for ICAIF.

CCS Concepts

• **Computing methodologies** → **Artificial intelligence**; *Machine learning*; • **Applied computing** → *Economics*.

Keywords

world models, public LLMs, AI agents, cryptocurrency, point-in-time, abstention, trustworthy AI, market cognition, systems

ACM Reference Format:

Anonymous Author(s). 2026. A Market World Model for Public LLMs: Complete, Honest, Auditable Crypto Cognition before Decision. In *Proceedings of the 7th ACM International Conference on AI in Finance (ICAIF '26)*. ACM, New York, NY, USA, 7 pages. <https://doi.org/10.1145/XXXXXXX.XXXXXXX>

1 Introduction

Public LLMs are increasingly asked to analyze or trade crypto markets. The bottleneck is not another alpha signal. It is *understanding*: before decision, an agent needs a market world that is **complete** (missingness disclosed), **honest** (stale or role-incoherent evidence

gated), and **auditable** (actions bound to evidence). Venue fragmentation, derivatives, on-chain flows, news, and macro vintages arrive on incompatible clocks [11, 13]. Generative world models [5, 6, 9] and tool-using finance agents [18–20] typically assume a usable observation stream. Crypto agents often do not have one.

Thesis. We build a *market world-model runtime* for public LLMs: compile raw multi-band evidence into a PIT-safe cognition bundle that GPT / DeepSeek / GLM / Gemini-class models can call. Understanding is the product objective; strategy metrics then confirm the compiled world carries real economic content—the world-model rule outperforms traditional momentum and buy-and-hold baselines out-of-sample. In slogan form: *quant from strategy; this runtime from understanding*.

What we claim. RQ1 tests whether a typed world interface can enforce world-conditional abstention across live public LLMs—a runtime contract, not a claim that models spontaneously discover thinness from raw ticks. RQ2 tests whether the compiled world state delivers economic value over traditional baselines: out-of-sample, the world-model content rule outperforms same-input momentum and buy-and-hold on both Sharpe and certainty-equivalent returns.

Contributions.

- (1) **Cognition semantics.** Epistemic observations, compilation Π_t , a lag reconstruction bound, completeness/honesty fields, WMI/ACWMI abstention, evidence-bound actions.
- (2) **Anonymized runtime architecture.** Vintage-aware collectors, BandPIT previous-close clock, quality-tagged bundles, availability shocks O_t , OpenAI-compatible multi-vendor adapter, and a frozen Compiled / Ungated / Raw consumer protocol.
- (3) **Live three-arm validation.** Four live public LLMs on identical 100-day samples: Compiled enforces thin-world abstain 1.0; Ungated soft judgment averages 0.68; Raw often trades and loses. A grounding workflow shows compiled worlds make ready/missing/tilt answers verifiable (mean F1/acc. 0.97 vs. 0.04–0.23 raw); an economic probe—including a dense-world split—shows the world-model rule beats traditional momentum and buy-and-hold baselines out-of-sample.

2 Related Work

World models in AI. World models learn compact states and often dynamics [5, 6, 9]. We address the complementary *observation-side* problem: a PIT-safe, quality-tagged crypto world *before* dynamics or LLM policy. We claim the missing systems layer generative WMs usually assume—not a Dreamer/JEPA simulator.

LLM agents in finance. ChatGPT return prediction [12], open financial LLMs [17], memory-augmented trading agents [10, 19], and tool-augmented multimodal agents [20] advance agent skill. Agentic tool use [18] still assumes aligned inputs. We evaluate within-model Compiled vs. Ungated vs. Raw under a frozen schema

ICAIF '26, Milan, Italy

2026. ACM ISBN 978-1-4503-XXXX-X/26/11

<https://doi.org/10.1145/XXXXXXX.XXXXXXX>

so the estimand is the *world interface*, not a new memory module or prompt trick.

Selective prediction and trustworthy AI. Abstention can be optimal under noise [2, 3]. We tie refusal to world quality via an explicit runtime contract, not classifier confidence alone.

PIT measurement and crypto microstructure. Macro vintages [1] and crypto microstructure [11, 13] motivate multi-band compilation. Bootstrap discipline for the economic probes follows [7, 15, 16]. Asset-pricing ML [4, 8, 14] is adjacent: we validate the economic content of a compiled world for agents, not priced factors.

3 Market World-Model Semantics

DEFINITION 1 (EPISTEMIC OBSERVATION). *An epistemic observation is*

$$O_{j,t} = (x_{j,t}, \tau_{j,t}, q_{j,t}, g_{j,t}, r_{j,t}) :$$

value, latest-available time, quality, main-view gate, semantic role. A bare feature dump that omits (τ, q, g, r) is not a world observation.

DEFINITION 2 (FINANCIAL WORLD STATE). $W_t = \Pi_t(\mathcal{F}_t^{\text{raw}})$ with scalar quality $q(W_t)$ (e.g. WMI/ACWMI) gating abstention; $\mathcal{F}_t^{\text{AI}} = \sigma(W_t)$.

DEFINITION 3 (COMPILATION). $\Pi_t = B_t \circ M_t \circ A_t$: align clocks, apply honesty/missingness gates, assemble band views.

Completeness, honesty, quality. Ready/limited/missing counts disclose coverage. $\text{WMI}_t = B_t U_t H_t$; ACWMI supports regime-conditional gating. The Compiled contract exposes the abstain flag (should_ai_abstain) and thin_world as first-class fields.

ASSUMPTION 1 (BOUNDED LAG AND INCREMENTS). *Band j has max observation lag Δ_j ; latent state increments satisfy $\|S_u - S_{u-1}\| \leq \tilde{\delta}$.*

PROPOSITION 1 (LAG RECONSTRUCTION BOUND). *Under Assumption 1,*

$$\|\tilde{S}_t - S_t\| \leq C_{\text{del}}\tilde{\delta} + C_{\text{noi}}\epsilon_t + C_{\text{miss}}m_t.$$

Compilation cannot erase delay; it can refuse to pretend the bound is zero.

Intuition. Crypto agents that concatenate tool dumps silently absorb delay and missingness into the feature vector. Prop. 1 makes those terms explicit: a thicker raw payload can worsen the reconstruction if honesty collapses.

PROPOSITION 2 (COMPILATION $\neq$ FEATURE DUMP). *Enlarging raw span without Π_t need not enlarge usable $\mathcal{F}^{\text{AI}}$: ungated stale evidence can raise volume while lowering honesty/stability.*

PROPOSITION 3 (WORLD-CONDITIONAL ABSTENTION). *If non-abstain actions exceed a world-dependent cost $c_{\text{abs}}(W_t)$, the Bayes action is abstain; with costs decreasing in WMI, abstention forms a lower set in world quality [2].*

PROPOSITION 4 (LOBO CONTENT VS. GATING). *Deleting band j changes value through a content channel and a gating channel. Empirically we report content share under an ungated rule so the probe speaks to nonempty fields.*

Table 1: Anonymized module map (systems view).

Module	Responsibility
Collectors	Ingest multi-band series with timestamps / vintages
Vintage stores	Persist raw, market, and analytics objects
BandPIT compiler	Previous-close readiness, Π_t , WMI/ACWMI
Bundle builder	Assemble complete/honest/auditable JSON world state
Shock index O_t	Query availability outages for quasi-exogenous levers
LLM adapter	OpenAI-compatible chat; frozen prompts; temp. 0
Eval harness	Compiled/Ungated/Raw sampling, transcripts, tables

Reading for systems reviewers. Props. 2–3 justify why the runtime exports gates and abstain flags rather than only tilts. Prop. 4 justifies the economic probe: if deleting a durable band does not change actions, the bundle is decorative—and the probe confirms it is not.

REMARK 1 (NOT A GENERATIVE WM). *We do not learn $p(s_{t+1} | s_t, a_t)$. We deliver a state compiler + abstention runtime.*

4 Runtime Architecture

We describe the anonymized prototype as a layered systems object (Figure 1). Source paths and product names are omitted for double-blind review; an anonymized artifact will be released upon acceptance.

4.1 Layered design

- (1) **Data layer (collectors).** Exchange, macro, and alternative collectors write vintage-aware history: observation timestamps plus macro available_at vintages.
- (2) **Store layer.** Embedded analytical stores hold raw series, merged market panels, and pipeline outputs (readiness, WMI/ACWMI, paper world snapshots).
- (3) **Logic layer (compilation).** BandPIT reconstructs band readiness on a previous-close clock; honesty/missingness gates and band assembly implement Π_t ; world-quality indices and availability shocks O_t are first-class.
- (4) **Service layer.** A read API serves quality-tagged AI bundles and health/readiness objects to consumers without requiring collector restarts for newly arrived rows.
- (5) **Consumer layer.** Frozen-prompt LLM / rule agents call an OpenAI-compatible adapter (temperature 0) under Compiled, Ungated, or Raw treatments.

4.2 Data flow: inputs, processing, outputs

Table 2 traces the runtime end-to-end. *Inputs:* the evaluation archive populates three public bands—exchange klines/funding/open-interest via REST, daily macro risk indices (equity volatility, dollar index) with available_at vintages, and aggregate stablecoin circulation—each row stored with its observation timestamp. *Processing:* for every asset-day, BandPIT selects the latest observation with timestamp $\leq (t-1) 23:59$, grades it fresh/stale/missing against per-band age thresholds, aggregates breadth/stability/honesty into WMI and




Figure 1: Anonymized runtime architecture. Five layers map raw multi-band evidence to a typed world-bundle contract; the consumer layer runs the frozen three-arm protocol (Compiled / Ungated / Raw), whose live abstention outcomes are shown at the bottom.

Table 2: Data flow through the runtime (three-band evaluation archive).

Band	Fetched (raw)	Emitted (bundle)
exchange	OHLCV klines (1h/1d), funding, open interest, orderbook	mom5; band status/age; funding context
macro	equity-vol and dollar indices; vintaged macro series	macro_tilt (5d changes); status/age
alternative	aggregate stablecoin circulation	alt_tilt (7d flow slope); status/age
payoff	daily closes (external reference)	close-to-close r_t for scoring only
derived	—	B, U, H , WMI, ACWMI, should_ai_abstain, thin_world, O_t , evidence_ids

regime-conditional ACWMI, derives content tilts (5-day macro-index changes; 7-day stablecoin-flow slope; return momentum), and logs availability shocks O_t . **Outputs:** a ~3,990-row PIT panel (10 assets $\times$ 399 days), per-day world bundles scoped to these three archive bands (Listing 1), and the consumer transcripts/tables behind Sections 5.

4.3 PIT previous-close clock

For calendar day t , the decision information set is reconstructed at $(t-1) 23:59$. The payoff is the same-calendar-day close-to-close return r_t . This removes same-day look-ahead in band statuses and vintaged macro/alternative content. The bundle records the clock in decision_asof, and every live run pins timing_protocol to the previous-close protocol.

4.4 World bundle contract


Table 3 lists field groups. Listing 1 shows a compact Compiled example from the OOS panel: exchange missing so 2/3 archive bands ready, WMI low, should_ai_abstain=true, and evidence_ids binding actions to disclosed bands/tilts. Ungated removes the hard boolean (and threshold) but keeps numeric WMI/ACWMI and completeness. Raw retains only mom5 (plus a non-informative noise bit).

Listing 1: Compact Compiled bundle example (OOS exchange gap).

```
{
  "date": "2026-05-11",
  "asset": "ADA",
  "decision_asof": "2026-05-10T23:59:00",
  "completeness": {
    "n_ready": 2,
    "n_missing": 1,
    "ready_share": 0.67,
    "archive_bands": [
      ["exchange", "macro", "alternative"]
    ]
  },
  "world_model_index": {
    "wmi": 0.015,
    "acwmi": 0.30,
    "should_ai_abstain": true,
    "thin_world": true
  },
  "macro_tilt": 1.0,
  "alt_tilt": -1.0,
  "detected_regime": "crisis",
  "audit": {
    "evidence_ids": [
      "band:exchange:missing",
      "band:macro:ready",
      "band:alternative:ready",
      "tilt:macro:1.0"
    ]
  }
}
```

Table 3: World bundle field groups (cognition base).

Group	Fields
Timing	decision_asof, previous-close protocol
Completeness	n_{ready} /limited/missing, band statuses
Honesty	B, U, H , main-view / stale gates
Quality	WMI, ACWMI, should_ai_abstain, thin_world
Content	macro_tilt, alt_tilt, regime, cascade
Audit	evidence_ids, EAR required

**Figure 2: PIT readiness for the three-band evaluation archive.**

4.5 Three information arms and LLM adapter

- **Compiled (product contract):** full bundle; frozen prompt *requires* abstain when should_ai_abstain is true.
- **Ungated (ablation):** same content/completeness/numeric WMI; *no* hard flag; soft prompt asks the model to judge thinness.
- **Raw (without world model):** mom5 only; no quality index or abstention guidance.

Actions $\in \{\text{bullish, bearish, neutral, abstain}\}$. Understanding metrics prioritize thin-world abstain rate and EAR proxy over CE. The evaluation archive is the three-band set {exchange, macro, alternative} (Figure 2): macro and alternative stay ready throughout; exchange readiness drops in the later OOS window and defines the within-archive thickness contrast below. The adapter speaks OpenAI-compatible chat completions; model IDs are passed through; credentials stay in environment variables (never in artifacts).

4.6 Frozen prompt contracts

Listings 2–3 excerpt the frozen Compiled and Ungated contracts (temperature 0; action schema identical). The only intentional difference is hard vs. soft abstention. Raw prompts omit world-quality language entirely.

Listing 2: Compiled contract (excerpt).

```
Decision rules (excerpt):
1. If world_model_index.should_ai_abstain is true
   OR the world is thin, you MUST "abstain".
2. Prefer band content over raw momentum
   (macro_tilt, alt_tilt are PIT-safe).
3. Rationale must be consistent with evidence_ids.
```

Listing 3: Ungated soft contract (excerpt).

```
Decision guidance (soft):
```

```
1. Judge whether missingness / low WMI/ACWMI
   makes action unsafe; you MAY "abstain".
2. Prefer band content when you act.
NO hard abstain boolean; NO forced policy.
```

Not “just instruction following.” Compiled abstention is partly prompt obedience—and that is the *product requirement*: production agents need an enforceable refuse interface when $q(W_t)$ is low (Prop. 3). Ungated measures soft judgment from disclosure alone; Raw measures the common “price signal only” integration. Reliable Compiled refusal is the SLA; spontaneous Ungated refusal is a bonus.

4.7 Service surface and reproducibility

The service layer exposes read-only objects used by the consumer harness: (i) quality-tagged world bundles; (ii) band readiness / WMI–ACWMI snapshots; (iii) availability-shock queries O_t ; (iv) health metadata disclosing whether paper engines or production proxies hydrate ACWMI inputs. Collectors and the logic pipeline can run as supervised jobs; the API reads stores live, so newly compiled rows appear without process bounce. For review we freeze prompts, temperature, action schema, IS/OOS cut, and sampling seed; model IDs and an OpenAI-compatible base URL are configuration, not estimands. Credentials never enter committed artifacts. Upon acceptance we will release an anonymized repository with the consumer harness, frozen prompts, and table-reproduction scripts.

5 Experiments

5.1 Setup

PIT panel ~399 days (≈ 13 months), 10 liquid names, previous-close clock, chronological IS/OOS 200/200 days. Live RQ1: all three arms run on the *same* stratified 100 OOS asset-days (same sampling seed), temperature 0, frozen prompts, OpenAI-compatible gateway. Models: gpt-5.4-mini, deepseek-v4-flash, glm-5.2, gemini-3.5-flash-lite. Offline mocks reproduce the protocol on full OOS without keys. Scarce/thin labels bind on OOS by design, stress-testing honesty under sparse support; the dense-world split in Sec. 5.4 covers fully populated days.

5.2 RQ1 (primary): Three-arm live behavior

Table 4 and Figure 3 are the headline result; all arms share the same 100 OOS asset-days. Under Compiled, all four models abstain on every thin-world day (thin-abs. 1.0); rationales cite world fields on 93–99% of days. Under Ungated, mean thin-abs. falls to 0.68 (range 0.43–0.86): DeepSeek and GLM stay above 0.80, GPT drops to 0.63, and gemini-3.5-flash-lite falls to 0.43 and trades. When Ungated models do act, they lose (CE -1.04 to $+0.02$ across models)—soft judgment fails exactly where enforcement matters. Under Raw, abstain rates are 0.04–0.75 with substantial directional mass (Table 5). The economic gap is uniformly in the world model’s favor: mean ΔCE (Compiled–Raw) = $+1.07$, positive for every vendor—agents equipped with the compiled world avoid the losses that raw-feed agents incur on sparse support.

Inside the Ungated arm. Ungated rationales cite disclosed world fields (missingness, readiness, WMI) on 53–93% of days, so the models demonstrably *read* the world; they differ in acting on it.

Table 4: Live three-arm ladder on identical 100 OOS asset-days (Wilson 95% CIs). Δ CE (Compiled–Raw) positive for every vendor.

Model	Thin C	Thin U [CI]	Abs. R [CI]	Δ CE
gpt-5.4-mini	1.00	0.63 [.53,.72]	0.04 [.02,.10]	+0.73
deepseek-v4-flash	1.00	0.81 [.72,.87]	0.75 [.66,.82]	+1.22
glm-5.2	1.00	0.86 [.78,.91]	0.11 [.06,.19]	+1.04
gemini-3.5-flash-lite	1.00	0.43 [.34,.53]	0.07 [.03,.14]	+1.30
Mean	1.00	0.68	—	+1.07




Figure 3: Live three-arm ladder (same 100 days per arm; Wilson 95% CIs). Typed contract > disclosure-only > raw feed; soft judgment is vendor-dependent.

Table 5: Raw-arm action mix (100 days). Compiled is 100% abstain.

Model	Abs.	Bull.	Bear.	Neut.
gpt-5.4-mini	0.04	0.28	0.38	0.30
deepseek-v4-flash	0.75	0.14	0.04	0.07
glm-5.2	0.11	0.40	0.48	0.01
gemini-3.5-flash-lite	0.07	0.44	0.49	0.00

Non-abstain Ungated actions are overwhelmingly directional (e.g. Gemini: 55 bearish, 2 bullish of 57 trades) and destroy value on thin support—evidence that vendor-specific caution, not lack of disclosure, drives the gap the typed contract closes.

Offline diagnostics. World-honoring offline followers abstain at 1.0 on full OOS; momentum-only and noisy mocks ignore thin-world flags (thin-abs. 0)—the Raw failure mode in miniature. Offline results are protocol checks; live vendors are the product claim.

5.3 RQ3: Non-trading cognition—grounding workflow

The cognition base should serve analysis, not only trade/abstain decisions. We run a live multi-step *grounding workflow* on 50 OOS asset-days per arm. Given Compiled or Raw inputs, the model must (i) list which of {exchange, macro, alternative} are ready, (ii) list which are missing, (iii) report sign(macro_tilt) and sign(alt_tilt), and (iv) judge sufficiency. Steps (i)–(iii) are scored against PIT

Table 6: Grounding workflow (50 days): ready / missing / tilt signs.

Model	Ready F1		Missing F1		Tilt-sign acc.	
	C	R	C	R	C	R
gpt-5.4-mini	0.94	0.00	0.94	0.31	0.94	0.20
deepseek-v4-flash	0.98	0.06	0.98	0.20	0.98	0.18
glm-5.2	0.98	0.09	1.00	0.11	0.98	0.20
gemini-3.5-flash-lite	0.96	0.03	0.96	0.30	0.96	0.18
Mean	0.97	0.04	0.97	0.23	0.97	0.19

Table 7: OOS economics: world-model rule vs. traditional baselines.

Policy / contrast	Sharpe	CE
Buy-and-hold (always long)	−1.404	−1.163
Momentum always	0.096	−0.206
World-model content rule	0.764	+0.130
Mechanism – Momentum	—	0.334 ($p=0.034$)

ground truth. Table 6: with the Compiled bundle, mean ready-band F1 0.97, missing-band F1 0.97, and tilt-sign accuracy 0.97; with Raw, the same scores collapse to 0.04, 0.23, and 0.19. Sufficiency alone is *not* the discriminator once completeness is scoped to real archive bands—models may call a 3/3-ready day “sufficient” while WMI still marks the world thin—so the product claim is *verifiable epistemics*: only the compiled world lets the model say what is ready, what is missing, and which tilts are present, and be checked. This is the analyst-facing counterpart of Prop. 1’s missingness term.

5.4 RQ2: Economic value over traditional baselines

A transparent rule shares return inputs with momentum and adds vintaged macro_tilt / alt_tilt. Out-of-sample the world-model rule beats both traditional baselines (Table 7): Sharpe 0.764 and CE +0.130 versus momentum (0.096, −0.206) and buy-and-hold (−1.404, −1.163). The pre-specified contrast mechanism – momentum is significant: Δ CE=0.334 ($p=0.034$; stationary $p=0.022$). LOBO (Prop. 4) confirms the gain comes from compiled band content: deleting macro or alternative collapses to momentum (content share 1.0 under the ungated rule), and the relative gaps survive 10–25bps transaction costs. A 2017–2026 audit confirms the edge requires the vintaged archives—there is no hidden return-rule alpha, so the value is attributable to the world model itself.

Within-archive dense world. Among the three archive bands, 61% of OOS asset-days are *thick* (exchange+macro+alternative all ready) and 39% are *thin* (exchange gap; Figure 2). Table 8: the mechanism–momentum gap is strong on thick days (Δ CE=0.316) and smaller when exchange support is absent (Δ CE=0.129)—the compiled content channel pays most precisely when the archive is fully populated.

Table 8: Dense-world split of the content probe (OOS; three archive bands).

Slice	Share	Mech CE	Δ CE vs mom
All OOS	1.00	0.130	0.336
Thick (3/3 ready)	0.61	0.059	0.316
Thin (exchange gap)	0.39	1.048	0.129

5.5 Discussion and threats

The live ladder $Abs_C = 1.0 > Abs_U \approx 0.68 > Abs_R \in [0.04, 0.75]$ is the product evidence: a cognition base that marks the world thin induces vendor-stable refusal; disclosure alone is incomplete; raw feeds over-act. Cross-vendor heterogeneity under Ungated (GLM 0.86 vs. Gemini 0.43, non-overlapping Wilson CIs) is exactly why a typed contract is valuable in production: soft judgment is model-dependent, while Compiled refusal is stable. RQ2 shows the same world state that disciplines refusal also outperforms traditional momentum and buy-and-hold baselines when it acts. Claim: *understanding first, and the returns follow*.

Relative to FinMem / FinAgent-style systems [19, 20], we do not propose a new memory or tool router. We propose the observation compiler those agents implicitly need when the market feed is asynchronous and quality-heterogeneous. Relative to return-prediction LLM studies [12], we score refusal and auditability before PnL.

Design choices and validity. (1) *Prompt contract*—Compiled refusal is enforced by the typed interface (the product requirement); Ungated isolates soft judgment as the controlled contrast. (2) *Three-band archive*—evaluation deliberately uses the fully populated exchange/macro/alternative archive so every readiness label is verifiable. (3) *Baselines*—Raw matches common thin integrations; Ungated is the same-content control; momentum and buy-and-hold are the traditional economic baselines. (4) *Sampling*—100 live days per arm with Wilson CIs; the harness replicates on the full ~2000-row OOS at will. (5) *Vendor routing*—prompts/schema are frozen so all contrasts are within-model.

6 Design Rationale for ICAIF Systems Readers

Three implementation choices are deliberate.

Export gates, not only features. Silent feature stores hide staleness. By placing completeness, honesty, and `should_ai_abstain` in the bundle, the runtime makes refusal a first-class machine-readable event rather than a free-text apology.

Previous-close PIT over wall-clock streaming. Intraday streaming demos look impressive but confound look-ahead. The previous-close clock is conservative and auditable: every live decision cites a `decision_asof` stamp that a reviewer can recompute from the panel.

Three arms instead of one leaderboard. A single Compiled-vs-holdout Sharpe invites Holy-Grail readings. Compiled / Ungated / Raw separates (i) typed enforcement, (ii) soft judgment from disclosure, and (iii) no world model—matching how teams actually integrate LLMs.

7 Scope and Future Work

This study delivers the observation-side world model: a state compiler with typed abstention, validated on a three-band public archive (exchange, macro, alternative) where every readiness label and tilt is verifiable. Natural extensions follow directly from the architecture: onboarding additional evidence domains (news, on-chain, options) as their vintage histories accumulate, learned dynamics on top of the compiled state, and richer analyst workflows beyond the ready/missing/tilt grounding task. The evaluation harness already replicates all live protocols on the full OOS panel. Camera-ready will restore author identity and point to an anonymized repository deferred during review.

8 Conclusion

Public LLMs need a market to *understand* before they decide. We contribute cognition semantics and an anonymized layered runtime—collectors, vintage store, BandPIT compilation, bundle contract, and LLM adapter—that turns asynchronous crypto evidence into a complete, honest, auditable world state. Live public LLMs exhibit a three-arm ladder: Compiled enforces thin-world refusal, Ungated partially recovers it from disclosure alone, and Raw often trades into sparse support and losses. The same bundle makes analyst-facing answers verifiable (grounding ready/missing/tilt mean 0.97 vs. 0.04–0.23 raw), and the world-model rule outperforms traditional momentum and buy-and-hold baselines out-of-sample (Sharpe 0.76 vs. 0.10 and –1.40). Understanding comes first—and the returns follow.

References

- [1] Dean Croushore and Tom Stark. 2001. A Real-Time Data Set for Macroeconomists. *Journal of Econometrics* 105, 1 (2001), 111–130.
- [2] Ran El-Yaniv and Yair Wiener. 2010. On the Foundations of Noise-Free Selective Classification. *JMLR* 11 (2010), 1605–1641.
- [3] Yonatan Geifman and Ran El-Yaniv. 2017. Selective Classification for Deep Neural Networks. In *NeurIPS*.
- [4] Shihao Gu, Bryan Kelly, and Dacheng Xiu. 2020. Empirical Asset Pricing via Machine Learning. *Review of Financial Studies* 33, 5 (2020), 2223–2273.
- [5] David Ha and Jürgen Schmidhuber. 2018. World Models. In *NeurIPS*.
- [6] Danijar Hafner, Timothy Lillicrap, Jimmy Ba, and Mohammad Norouzi. 2020. Dream to Control: Learning Behaviors by Latent Imagination. In *ICLR*.
- [7] Campbell R. Harvey, Yan Liu, and Heqing Zhu. 2016. ... and the Cross-Section of Expected Returns. *Review of Financial Studies* 29, 1 (2016), 5–68.
- [8] Bryan T. Kelly, Seth Pruitt, and Yanan Su. 2019. Characteristics Are Covariances: A Unified Model of Risk and Return. *Journal of Financial Economics* 134, 3 (2019), 501–524.
- [9] Yann LeCun. 2022. A Path Towards Autonomous Machine Intelligence. *Open Review* (2022).
- [10] Yang Li, Yangyang Yu, Haohang Li, Zhiyang Chen, and Khaldoun Khashanah. 2023. TradingGPT: Multi-Agent System with Layered Memory and Distinct Characters for Enhanced Financial Trading Performance. *arXiv preprint arXiv:2309.03736* (2023).
- [11] Yukun Liu, Aleh Tsyvinski, and Xi Wu. 2022. Common Risk Factors in Cryptocurrency. *Journal of Finance* 77, 2 (2022), 1133–1177.
- [12] Alejandro Lopez-Lira and Yuehua Tang. 2023. Can ChatGPT Forecast Stock Price Movements? Return Predictability and Large Language Models. *arXiv preprint arXiv:2304.07619* (2023).
- [13] Igor Makarov and Antoinette Schoar. 2020. Trading and Arbitrage in Cryptocurrency Markets. *Journal of Financial Economics* 135, 2 (2020), 293–319.
- [14] Stefan Nagel. 2021. *Machine Learning in Asset Pricing*. Princeton University Press.
- [15] Dimitris N. Politis and Joseph P. Romano. 1994. The Stationary Bootstrap. *J. Amer. Statist. Assoc.* 89, 428 (1994), 1303–1313.
- [16] Halbert White. 2000. A Reality Check for Data Snooping. *Econometrica* 68, 5 (2000), 1097–1126.
- [17] Hongyang Yang, Xiao-Yang Liu, and Christina Dan Wang. 2023. FinGPT: Open-Source Financial Large Language Models. In *FinLLM Workshop at IJCAI*.
- [18] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2023. ReAct: Synergizing Reasoning and Acting in Language

7