

A Market World Model for Public LLMs: A Typed Observation Contract before Decision

Anonymous Author(s)

Abstract

Public LLMs asked to analyze or trade crypto markets fail before strategy: they lack a complete, honest, auditable *market observation interface*. We present an anonymized *world-model runtime* that compiles asynchronous multi-band evidence into a point-in-time (PIT) world bundle with a typed abstention contract. Epistemic observations $O_{j,t} = (x, \tau, q, g, r)$ and compilation Π_t yield $\mathcal{F}_t^{AI}$; completeness, honesty (B, U, H) , and WMI/ACWMI expose machine-readable refusal; evidence IDs make actions auditable. The primary claim is *trustworthy interfacing*, not a new trading agent: on identical 100 OOS asset-days, four live public LLMs (temperature 0) form a behavior ladder—asked *blind*, every model refuses (1.0); given a momentum fragment they over-trade (abstain 0.04–0.75) and lose; given the compiled world with a hard flag they abstain on every thin day (1.0); Ungated (same content, no hard flag) averages only 0.68. A full-OOS checkpointed sweep ($n=2000$) keeps Compiled at 1.000 for all vendors. Even on archive-thick days (all three evaluation bands ready), Raw still over-acts and often loses while Compiled continues to refuse—readiness counts alone do not replace the typed contract. A grounding workflow recovers ready/missing/tilt answers at mean F1/acc. 0.97 from Compiled vs. 0.04–0.23 from Raw. As a secondary *content* probe (not an LLM leaderboard), a transparent world-model rule that reads the same vintaged tilts beats momentum and buy-and-hold out-of-sample (Sharpe 0.76 vs. 0.10 / -1.40 ; $\Delta CE=0.334$, $p=0.034$). We detail the runtime (collectors, vintage store, BandPIT, bundle schema, LLM adapter) as an ICAIF systems contribution.

CCS Concepts

• **Computing methodologies** → **Artificial intelligence**; *Machine learning*; • **Applied computing** → *Economics*.

Keywords

world models, public LLMs, AI agents, cryptocurrency, point-in-time, abstention, trustworthy AI, observation contract, systems

ACM Reference Format:

Anonymous Author(s). 2026. A Market World Model for Public LLMs: A Typed Observation Contract before Decision. In *Proceedings of the 7th ACM International Conference on AI in Finance (ICAIF '26)*. ACM, New York, NY, USA, 7 pages. <https://doi.org/10.1145/XXXXXXX.XXXXXXX>

1 Introduction

Public LLMs are increasingly asked to analyze or trade crypto markets. The bottleneck is not another alpha signal. It is *understanding*: before decision, an agent needs a market world that is **complete** (missingness disclosed), **honest** (stale or role-incoherent evidence

gated), and **auditable** (actions bound to evidence). Venue fragmentation, derivatives, on-chain flows, news, and macro vintages arrive on incompatible clocks [11, 13]. Generative world models [5, 6, 9] and tool-using finance agents [18–20] typically assume a usable observation stream. Crypto agents often do not have one.

Thesis. We build a market *world-model runtime* for public LLMs: compile raw multi-band evidence into a PIT-safe bundle that GPT / DeepSeek / GLM / Gemini-class models can call under a typed refuse contract. The product objective is *trustworthy observation*, not a new portfolio policy. Strategy metrics appear only as a secondary probe that the compiled fields are nonempty. In slogan form: *quant from strategy; this runtime from understanding—enforced at the interface*.

What we claim. RQ1 tests whether a typed world interface can *enforce* world-conditional abstention across live public LLMs—a runtime contract, not a claim that models spontaneously discover thinness from raw ticks. RQ2 tests whether the compiled world makes analyst-facing epistemic answers verifiable. RQ3 is a *content* probe: when a transparent rule reads the same vintaged tilts, does the compiled world beat traditional momentum and buy-and-hold? We do *not* claim an LLM trading leaderboard.

Contributions.

- (1) **Cognition semantics.** Epistemic observations, compilation Π_t , a lag reconstruction bound, completeness/honesty fields, WMI/ACWMI abstention, evidence-bound actions.
- (2) **Anonymized runtime architecture.** Vintage-aware collectors, BandPIT previous-close clock, quality-tagged bundles, availability shocks O_t , OpenAI-compatible multi-vendor adapter, and a frozen Compiled / Ungated / Raw / Blind consumer protocol.
- (3) **Live interface validation.** Four-arm ladder on identical days: Blind refuses, Raw over-trades and loses, Ungated is vendor-dependent, Compiled enforces thin-world abstain 1.0 (replicated on full OOS for all vendors). A band-thick split shows the contract still binds when all archive bands are ready; grounding recovers ready/missing/tilt at 0.97 vs. 0.04–0.23 raw; a secondary content rule confirms the tilts are economically nonempty.

2 Related Work

World models in AI. World models learn compact states and often dynamics [5, 6, 9]. We address the complementary *observation-side* problem: a PIT-safe, quality-tagged crypto world *before* dynamics or LLM policy. We claim the missing systems layer generative WMs usually assume—not a Dreamer/JEPA simulator.

LLM agents in finance. ChatGPT return prediction [12], open financial LLMs [17], memory-augmented trading agents [10, 19], and tool-augmented agents [18, 20] advance agent skill and orchestration. Recent ICAIF-facing lines further emphasize agent

ICAIF '26, Milan, Italy

2026. ACM ISBN 978-1-4503-XXXX-X/26/11

<https://doi.org/10.1145/XXXXXXX.XXXXXXX>

benchmarks, retrieval, and research-judge protocols. We deliberately do *not* propose another memory router or debate society: we evaluate within-model Compiled vs. Ungated vs. Raw vs. Blind under a frozen schema so the estimand is the *observation contract* those agents still assume.

Selective prediction and trustworthy AI. Abstention can be optimal under noise [2, 3]. We tie refusal to world quality via an explicit runtime contract, not classifier confidence alone.

PIT measurement and crypto microstructure. Macro vintages [1] and crypto microstructure [11, 13] motivate multi-band compilation. Bootstrap discipline for the economic probes follows [7, 15, 16]. Asset-pricing ML [4, 8, 14] is adjacent: we validate the economic content of a compiled world for agents, not priced factors.

3 Market World-Model Semantics

Definition 3.1 (Epistemic observation). An epistemic observation is

$$O_{j,t} = (x_{j,t}, \tau_{j,t}, q_{j,t}, g_{j,t}, r_{j,t}) :$$

value, latest-available time, quality, main-view gate, semantic role. A bare feature dump that omits (τ, q, g, r) is not a world observation.

Definition 3.2 (Financial world state). $W_t = \Pi_t(\mathcal{F}_t^{\text{raw}})$ with scalar quality $q(W_t)$ (e.g. WMI/ACWMI) gating abstention; $\mathcal{F}_t^{\text{AI}} = \sigma(W_t)$.

Definition 3.3 (Compilation). $\Pi_t = B_t \circ M_t \circ A_t$: align clocks, apply honesty/missingness gates, assemble band views.

Completeness, honesty, quality. Ready/limited/missing counts disclose coverage. $\text{WMI}_t = B_t U_t H_t$; ACWMI supports regime-conditional gating. The Compiled contract exposes the abstain flag (should_ai_abstain) and thin_world as first-class fields.

Assumption 3.4 (Bounded lag and increments). Band j has max observation lag Δ_j ; latent state increments satisfy $\|S_u - S_{u-1}\| \leq \bar{\delta}$.

PROPOSITION 3.5 (LAG RECONSTRUCTION BOUND). *Under Assumption 1,*

$$\|\tilde{S}_t - S_t\| \leq C_{\text{del}}\bar{\delta} + C_{\text{noi}}\varepsilon_t + C_{\text{miss}}m_t.$$

Compilation cannot erase delay; it can refuse to pretend the bound is zero.

Intuition. Crypto agents that concatenate tool dumps silently absorb delay and missingness into the feature vector. Prop. 3.5 makes those terms explicit: a thicker raw payload can worsen the reconstruction if honesty collapses.

PROPOSITION 3.6 (COMPILATION $\neq$ FEATURE DUMP). *Enlarging raw span without Π_t need not enlarge usable $\mathcal{F}^{\text{AI}}$: ungated stale evidence can raise volume while lowering honesty/stability.*

PROPOSITION 3.7 (WORLD-CONDITIONAL ABSTENTION). *If non-abstain actions exceed a world-dependent cost $c_{\text{abs}}(W_t)$, the Bayes action is abstain; with costs decreasing in WMI, abstention forms a lower set in world quality [2].*

PROPOSITION 3.8 (LOBO CONTENT VS. GATING). *Deleting band j changes value through a content channel and a gating channel. Empirically we report content share under an ungated rule so the probe speaks to nonempty fields.*

Table 1: Anonymized module map (systems view).

Module	Responsibility
Collectors	Ingest multi-band series with timestamps / vintages
Vintage stores	Persist raw, market, and analytics objects
BandPIT compiler	Previous-close readiness, Π_t , WMI/ACWMI
Bundle builder	Assemble complete/honest/auditable JSON world state
Shock index O_t	Query availability outages for quasi-exogenous levers
LLM adapter	OpenAI-compatible chat; frozen prompts; temp. 0
Eval harness	Compiled/Ungated/Raw/Blind sampling, transcripts, tables

Reading for systems reviewers. Props. 3.6–3.7 justify why the runtime exports gates and abstain flags rather than only tilts. Prop. 3.8 justifies the economic probe: if deleting a durable band does not change actions, the bundle is decorative—and the probe confirms it is not.

REMARK 3.9 (NOT A GENERATIVE WM). *We do not learn $p(s_{t+1} | s_t, a_t)$. We deliver a state compiler + abstention runtime.*

4 Runtime Architecture

We describe the anonymized prototype as a layered systems object (Figure 1; module map in Table 1). Source paths and product names are omitted for double-blind review; an anonymized artifact will be released upon acceptance.

4.1 Layered design

- (1) **Data layer (collectors).** Exchange, macro, and alternative collectors write vintage-aware history: observation timestamps plus macro available_at vintages.
- (2) **Store layer.** Embedded analytical stores hold raw series, merged market panels, and pipeline outputs (readiness, WMI/ACWMI, paper world snapshots).
- (3) **Logic layer (compilation).** BandPIT reconstructs band readiness on a previous-close clock; honesty/missingness gates and band assembly implement Π_t ; world-quality indices and availability shocks O_t are first-class.
- (4) **Service layer.** A read API serves quality-tagged AI bundles and health/readiness objects to consumers without requiring collector restarts for newly arrived rows.
- (5) **Consumer layer.** Frozen-prompt LLM / rule agents call an OpenAI-compatible adapter (temperature 0) under Compiled, Ungated, or Raw treatments.

4.2 Data flow: inputs, processing, outputs

Table 2 traces the runtime end-to-end. *Inputs:* the evaluation archive populates three public bands—exchange klines/funding/open-interest via REST, daily macro risk indices (equity volatility, dollar index) with available_at vintages, and aggregate stablecoin circulation—each row stored with its observation timestamp. *Processing:* for




Figure 1: Anonymized runtime architecture. Five layers map raw multi-band evidence to a typed world-bundle contract; the consumer layer runs the frozen four-arm protocol (Compiled / Ungated / Raw / Blind), whose live abstention outcomes are shown at the bottom.

every asset-day, BandPIT selects the latest observation with timestamp $\leq (t-1) 23:59$, grades it fresh/stale/missing against per-band age thresholds, aggregates breadth/stability/honesty into WMI and regime-conditional ACWMI, derives content tilts (5-day macro-index changes; 7-day stablecoin-flow slope; return momentum), and logs availability shocks O_t . *Outputs:* a ~3,990-row PIT panel (10 assets $\times$ 399 days), per-day world bundles scoped to these three archive bands (Listing 1), and the consumer transcripts/tables behind Section 5.

4.3 PIT previous-close clock

For calendar day t , the decision information set is reconstructed at $(t-1) 23:59$. The payoff is the same-calendar-day close-to-close return r_t . This removes same-day look-ahead in band statuses and vintaged macro/alternative content. The bundle records the clock in decision_asof, and every live run pins timing_protocol to the previous-close protocol.

4.4 World bundle contract

Table 3 lists field groups. Listing 1 shows a compact Compiled example from the OOS panel: exchange missing so 2/3 archive bands ready, WMI low, should_ai_abstain=true, and evidence_ids binding actions to disclosed bands/tilts. Ungated removes the hard boolean (and threshold) but keeps numeric WMI/ACWMI and completeness. Raw retains only mom5 (plus a non-informative noise bit).

Table 2: Data flow through the runtime (three-band evaluation archive).


Band	Fetched (raw)	Emitted (bundle)
exchange	OHLCV klines (1h/1d), funding, open interest, orderbook	mom5; band status/age; funding context
macro	equity-vol and dollar indices; vintaged macro series	macro_tilt (5d changes); status/age
alternative	aggregate stablecoin circulation	alt_tilt (7d flow slope); status/age
payoff	daily closes (external reference)	close-to-close r_t for scoring only
derived	—	B, U, H , WMI, ACWMI, should_ai_abstain, thin_world, O_t , evidence_ids

Listing 1: Compact Compiled bundle example (OOS exchange gap).

```
{
  "date": "2026-05-11",
  "asset": "ADA",
  "decision_asof": "2026-05-10T23:59:00",
  "completeness": {
    "n_ready": 2,
    "n_missing": 1,
    "ready_share": 0.67,
    "archive_bands": [
      ["exchange", "macro", "alternative"]
    ]
  },
  "world_model_index": {
    "wmi": 0.015,
    "acwmi": 0.30,
    "should_ai_abstain": true,
    "thin_world": true
  },
  "macro_tilt": 1.0,
  "alt_tilt": -1.0,
  "detected_regime": "crisis"
}
```

Table 3: World bundle field groups (cognition base).

Group	Fields
Timing	decision_asof, previous-close protocol
Completeness	n_{ready} /limited/missing, band statuses
Honesty	B, U, H , main-view / stale gates
Quality	WMI, ACWMI, should_ai_abstain, thin_world
Content	macro_tilt, alt_tilt, regime, cascade
Audit	evidence_ids, EAR required

**Figure 2: PIT readiness for the three-band evaluation archive.**

```
"audit":{"evidence_ids":[
  "band:exchange:missing", "band:macro:ready",
  "band:alternative:ready", "tilt:macro:1.0"]}}
```

4.5 Four information arms and LLM adapter

- **Compiled (product contract):** full bundle; frozen prompt *requires* abstain when should_ai_abstain is true.
- **Ungated (strong content control):** same content/completeness/numeric WMI; *no* hard flag; soft prompt asks the model to judge thinness. This is the non-strawman ablation: disclosure without enforcement.
- **Raw (common thin integration):** mom5 only; no quality index or abstention guidance—matching the widespread “price signal in, action out” wiring, not a cartoon baseline.
- **Blind (direct ask):** the retail question verbatim—“Today is d . How should I trade a ?”—with *no* data feed at all.

Actions $\in \{\text{bullish, bearish, neutral, abstain}\}$. Primary metrics are thin-world abstain rate and EAR proxy; CE is reported as a safety side-effect (avoided loss), not as an alpha claim. The evaluation archive is the three-band set {exchange, macro, alternative} (Figure 2): macro and alternative stay ready throughout; exchange readiness drops in the later OOS window and defines the within-archive *band-thick* contrast below. In this scarce panel WMI stays below the refuse threshold on every asset-day, so Compiled abstention is the intended stress test of the contract under sparse support. The adapter speaks OpenAI-compatible chat completions; credentials stay in environment variables (never in artifacts).

4.6 Frozen prompt contracts

Listings 2–3 excerpt the frozen Compiled and Ungated contracts (temperature 0; action schema identical). The only intentional difference is hard vs. soft abstention. Raw prompts omit world-quality language entirely.

Listing 2: Compiled contract (excerpt).

```
Decision rules (excerpt):
1. If world_model_index.should_ai_abstain is true
   OR the world is thin, you MUST "abstain".
2. Prefer band content over raw momentum
   (macro_tilt, alt_tilt are PIT-safe).
3. Rationale must be consistent with evidence_ids.
```

Listing 3: Ungated soft contract (excerpt).

```
Decision guidance (soft):
1. Judge whether missingness / low WMI/ACWMI
   makes action unsafe; you MAY "abstain".
2. Prefer band content when you act.
NO hard abstain boolean; NO forced policy.
```

Not “just instruction following.” Compiled abstention is partly prompt obedience—and that is the *product requirement*: production agents need an enforceable refuse interface when $q(W_t)$ is low (Prop. 3.7). Ungated measures soft judgment from disclosure alone; Raw measures the common “price signal only” integration. Reliable Compiled refusal is the SLA; spontaneous Ungated refusal is a bonus.

4.7 Service surface and reproducibility

The service layer exposes read-only objects used by the consumer harness: (i) quality-tagged world bundles; (ii) band readiness / WMI–ACWMI snapshots; (iii) availability-shock queries O_t ; (iv) health metadata disclosing whether paper engines or production proxies hydrate ACWMI inputs. Collectors and the logic pipeline can run as supervised jobs; the API reads stores live, so newly compiled rows appear without process bounce. For review we freeze prompts, temperature, action schema, IS/OOS cut, and sampling seed; model IDs and an OpenAI-compatible base URL are configuration, not estimands. Credentials never enter committed artifacts. Upon acceptance we will release an anonymized repository with the consumer harness, frozen prompts, and table-reproduction scripts.

5 Experiments

5.1 Setup

PIT panel ~399 days, 10 liquid names, previous-close clock, chronological IS/OOS 200/200 days ($\approx 2,000$ OOS asset-days). Live RQ1 uses the *same* stratified 100 OOS asset-days across arms (seed fixed), temperature 0, frozen prompts, and an OpenAI-compatible gateway; a checkpointed full-OOS sweep tests population scale. Vendor models are GPT-, DeepSeek-, GLM-, and Gemini-class flash/mini IDs listed in Table 4. Scarce/thin labels bind on OOS by design; Sec. 5.4 covers fully populated days.

Table 4: Live abstention ladder on identical 100 OOS asset-days (Wilson 95% CIs); Blind abstains 1.0 for every model. Δ CE is avoided-loss, not an alpha claim.

Model	Thin C	Thin U [CI]	Abs. R [CI]	Δ CE
gpt-5.4-mini	1.00	0.63 [.53,.72]	0.04 [.02,.10]	+0.73
deepseek-v4-flash	1.00	0.81 [.72,.87]	0.75 [.66,.82]	+1.22
glm-5.2	1.00	0.86 [.78,.91]	0.11 [.06,.19]	+1.04
gemini-3.5-flash-lite	1.00	0.43 [.34,.53]	0.07 [.03,.14]	+1.30
Mean	1.00	0.68	—	+1.07

Live four-arm ladder: blind refuses, raw fragment over-acts, compiled acts on world quality (Wilson 95% CIs, $n = 100$)




Figure 3: Live four-arm interface ladder (same 100 days; Wilson 95% CIs). Blind refuses; Raw over-trades; Ungated is vendor-dependent; Compiled enforces the typed refuse contract.

Table 5: Raw-arm action mix (100 days). Compiled abstains on all thin days; Blind (direct ask, no feed) abstains 100% for every model.

Model	Abs.	Bull.	Bear.	Neut.
gpt-5.4-mini	0.04	0.28	0.38	0.30
deepseek-v4-flash	0.75	0.14	0.04	0.07
glm-5.2	0.11	0.40	0.48	0.01
gemini-3.5-flash-lite	0.07	0.44	0.49	0.00

5.2 RQ1 (primary): Live interface behavior across information arms

Table 4 and Figure 3 are the headline result; all arms share the same 100 OOS asset-days. Under Compiled, all four models abstain on every thin-world day (thin-abs. 1.0); rationales cite world fields on 93–99% of days—this is prompt-contract obedience by design (the product SLA), not a claim of spontaneous thinness discovery. Under Ungated, mean thin-abs. falls to 0.68 (range 0.43–0.86): DeepSeek and GLM stay above 0.80, GPT drops to 0.63, and gemini-3.5-flash-lite falls to 0.43 and trades. When Ungated models do act on this 100-day draw, they lose (CE -1.04 to $+0.02$)—soft judgment is unstable exactly where enforcement matters. Under Raw, abstain rates are 0.04–0.75 with substantial directional mass (Table 5). Safety side-effect: mean Δ CE (Compiled–Raw) = $+1.07$, positive for every vendor, because Compiled sits out the losses Raw incurs on sparse support.

Table 6: Full-OOS band-thick split (3/3 ready; WMI still thin). Compiled refuses; Raw over-acts.

Model	Abs. U	Abs. R	CE R
gpt-5.4-mini	0.58	0.00	−0.28
deepseek-v4-flash	0.42	0.75	+0.04
glm-5.2 ($n=398$ R)	0.58	0.09	−0.87
gemini-3.5-flash-lite	0.34	0.00	−0.25

Full-OOS replication. To rule out subsample luck we checkpointed the three data arms on the *full* OOS panel ($n=2000$ asset-days; every OOS day is thin-world under the scarce-support design). Compiled thin-abstain is 1.000 [.998, 1.000] for *all four* vendors—the typed refuse contract is a population property, not a lucky draw. Ungated rates tighten the Wilson intervals and keep the soft-judgment ladder, with mean thin-abs. 0.56 (GPT 0.666 [.645, .686]; GLM 0.679 [.654, .703] on $n=1402$ completed calls; DeepSeek 0.540 [.518, .561]; Gemini 0.359 [.338, .380]). Vendor *ranking* within Ungated shifts relative to the 100-day sample (DeepSeek is less cautious at scale), but no vendor matches Compiled. Raw fragment over-trading remains decisive where the scrap is trusted: GPT and Gemini abstain at 0.000 [.000, .002]; DeepSeek stays cautious at 0.748 [.728, .767] ($n=1925$). Because the full OOS window is entirely thin, Compiled CE is identically zero by construction—thick-day economics are the dense-world probe of Sec. 5.4, not this honesty stress test.

The Blind arm: what happens if you just ask. The product-realistic control is the question a retail user actually types: “*Today is d. How should I trade BTC?*” with no data feed. Run live on the same 100 asset-days, all four models abstain on 100% of days (CE = 0; zero directional actions). A bare public LLM is not a market agent—it refuses. The failure mode that costs money is the middle case: give the same model a *fragment* (Raw) and it starts trading (abstain as low as 0.04) and loses. Compiled matches Blind’s refusal on thin days *by contract*, while still exposing the structured world that RQ2 shows is analytically usable and that RQ3 shows is economically nonempty when consumed by a transparent rule.

Half-B patch: band-thick days still need the contract. Archive readiness alone is not a refuse policy. On full-OOS *band-thick* days (exchange+macro+alternative all ready; 61% of OOS, $n \approx 1224$), WMI still marks every day thin under the scarce-support design, so Compiled abstains at 1.0. Ungated soft judgment only partially recovers (mean abs. ≈ 0.48 ; Gemini 0.34). Raw remains the dangerous middle: GPT and Gemini abstain at 0.0 with CE -0.28 and -0.25 . Table 6: when the archive *looks* complete, fragment feeds still over-act and lose, while the typed contract continues to refuse—exactly the production failure mode a readiness dashboard would miss.

Inside the Ungated arm. Ungated rationales cite disclosed world fields on 53–93% of days, so models *read* the world; they differ in acting on it. Vendor-specific caution—not missing disclosure—drives the gap the typed contract closes.

Offline diagnostics. World-honoring offline followers abstain at 1.0 on full OOS; momentum-only mocks ignore thin-world flags (thin-abs. 0). Offline results are protocol checks; live vendors are the product claim.

Table 7: Grounding workflow (50 days): ready / missing / tilt signs.

Model	Ready F1		Missing F1		Tilt-sign acc.	
	C	R	C	R	C	R
gpt-5.4-mini	0.94	0.00	0.94	0.31	0.94	0.20
deepseek-v4-flash	0.98	0.06	0.98	0.20	0.98	0.18
glm-5.2	0.98	0.09	1.00	0.11	0.98	0.20
gemini-3.5-flash-lite	0.96	0.03	0.96	0.30	0.96	0.18
Mean	0.97	0.04	0.97	0.23	0.97	0.19

5.3 RQ2: Non-trading cognition—grounding workflow

The cognition base should serve analysis, not only trade/abstain decisions. We run a live multi-step *grounding workflow* on 50 OOS asset-days per arm. Given Compiled or Raw inputs, the model must (i) list which of {exchange, macro, alternative} are ready, (ii) list which are missing, (iii) report $\text{sign}(\text{macro_tilt})$ and $\text{sign}(\text{alt_tilt})$, and (iv) judge sufficiency. Steps (i)–(iii) are scored against PIT ground truth. Table 7: with the Compiled bundle, mean ready-band F1 0.97, missing-band F1 0.97, and tilt-sign accuracy 0.97; with Raw, the same scores collapse to 0.04, 0.23, and 0.19. Sufficiency alone is *not* the discriminator once completeness is scoped to real archive bands—models may call a 3/3-ready day “sufficient” while WMI still marks the world thin—so the product claim is *verifiable epistemics*: only the compiled world lets the model say what is ready, what is missing, and which tilts are present, and be checked. This is the analyst-facing counterpart of Prop. 3.5’s missingness term.

5.4 RQ3: Secondary content probe (not an LLM leaderboard)

RQ1–RQ2 establish the interface. RQ3 asks a narrower question: are the compiled tilts economically nonempty, or decorative? A transparent rule shares return inputs with momentum and adds vintaged macro_tilt / alt_tilt —no LLM in the loop. Out-of-sample (199 days $\times$ 10 assets) this content rule beats traditional baselines (Table 8, Figure 4): Sharpe 0.764 / CE +0.130 versus momentum (0.096, −0.206) and buy-and-hold (−1.404, −1.163). The pre-specified contrast mechanism – momentum is significant under block bootstrap ($\Delta\text{CE}=0.334$, $p=0.034$). ΔCE is positive for all ten assets (Figure 4, right). LOBO (Prop. 3.8) attributes the gain to band content (deleting macro/alternative collapses to momentum); gaps survive 10–25bps costs. A 2017–2026 audit with non-vintaged proxy tilts shows no edge ($\Delta\text{CE} \approx 0$, $p=0.81$): value appears where the vintage-compiled world exists. This supports “the bundle is not empty,” not “our LLM is a trading SOTA.”


Within-archive dense world. Among the three archive bands, 61% of OOS asset-days are band-thick (3/3 ready) and 39% are band-thin (exchange gap; Figure 2). Table 9: the mechanism–momentum gap is stronger on band-thick days ($\Delta\text{CE}=0.316$) than on exchange-gap days (0.129)—the content channel pays most when the archive is populated, complementary to RQ1’s refuse-under-thin stress test.

Table 8: Secondary content probe: world-model rule vs. traditional baselines (no LLM).

Policy / contrast	Sharpe	CE
Buy-and-hold (always long)	−1.404	−1.163
Momentum always	0.096	−0.206
World-model content rule	0.764	+0.130
Mechanism – Momentum	—	0.334 ($p=0.034$)

Table 9: Dense-world split of the content probe (OOS; three archive bands).

Slice	Share	Mech CE	ΔCE vs mom
All OOS	1.00	0.130	0.336
Thick (3/3 ready)	0.61	0.059	0.316
Thin (exchange gap)	0.39	1.048	0.129

**Figure 4: Secondary content probe (no LLM). Left: equity curves; right: ΔCE vs. momentum positive for every asset.**

5.5 Discussion and threats

The product evidence is the interface ladder: Blind refuses (no product), Raw over-acts and loses, Ungated is vendor-dependent (≈ 0.68 ; GLM 0.86 vs. Gemini 0.43), Compiled enforces thin-world abstain 1.0—including on band-thick days where a readiness dashboard would look green. Claim: *understanding first, enforced at the contract; content economics are secondary*.

Relative to FinMem / FinAgent-style systems [19, 20], we do not propose a new memory or tool router. We propose the observation compiler and refuse SLA those agents need when feeds are asynchronous and quality-heterogeneous. Relative to return-prediction LLM studies [12], we score refusal and auditability before PnL.

Design choices and validity. (1) *Prompt contract*—Compiled refusal is the product SLA; Ungated is the strong same-content control. (2) *Three-band archive*—every readiness label is verifiable; band-thick vs. band-thin separates archive coverage from WMI thinness. (3) *Baselines*—Raw matches common mom-only wiring; RQ3 baselines are momentum/buy-and-hold for the content probe only. (4) *Sampling*—100-day headline ladder plus full-OOS checkpoints ($n=2000$; GLM Raw incomplete). (5) *What we do not claim*—LLM trading SOTA, generative dynamics, or spontaneous thinness discovery without the boolean.

6 Design Rationale for ICAIF Systems Readers

Three implementation choices are deliberate.

Export gates, not only features. Silent feature stores hide staleness. By placing completeness, honesty, and `should_ai_abstain` in the bundle, the runtime makes refusal a first-class machine-readable event rather than a free-text apology.

Previous-close PIT over wall-clock streaming. Intraday streaming demos look impressive but confound look-ahead. The previous-close clock is conservative and auditable: every live decision cites a `decision_asof` stamp that a reviewer can recompute from the panel.

Four arms instead of one leaderboard. A single Compiled-vs-holdout Sharpe invites Holy-Grail readings. Compiled / Ungated / Raw / Blind separates (i) typed enforcement, (ii) soft judgment from disclosure, (iii) common thin integrations, and (iv) no feed—matching how people actually query LLMs about markets, without turning the paper into a trading contest.

7 Scope and Future Work

This study delivers an observation-side world-model *interface*: a state compiler with typed abstention on a three-band public archive. We do not claim LLM portfolio SOTA, generative dynamics, or coverage of news/on-chain/options until those vintages exist. Natural extensions: additional evidence domains, learned dynamics on the compiled state, richer analyst workflows, and optional thick-regime live trading studies once WMI leaves the scarce floor. Camera-ready will restore author identity and release an anonymized harness (frozen prompts, sampling seed, table-reproduction scripts) deferred during review.

8 Conclusion

Public LLMs need a market they can *safely observe* before they decide. We contribute cognition semantics and an anonymized layered runtime that turns asynchronous crypto evidence into a complete, honest, auditable world bundle with a typed refuse contract. Live public LLMs exhibit a four-arm interface ladder—Blind refuses, Raw over-trades and loses, Ungated is vendor-dependent, Compiled enforces thin-world abstain 1.0 even on band-thick days—and the same bundle makes analyst answers verifiable (grounding 0.97 vs. 0.04–0.23). A secondary content rule shows the tilts are economically nonempty versus momentum and buy-and-hold. Understanding comes first—enforced at the interface.

References

- [1] Dean Croushore and Tom Stark. 2001. A Real-Time Data Set for Macroeconomists. *Journal of Econometrics* 105, 1 (2001), 111–130.
- [2] Ran El-Yaniv and Yair Wiener. 2010. On the Foundations of Noise-Free Selective Classification. *Journal of Machine Learning Research* 11 (2010), 1605–1641.
- [3] Yonatan Geifman and Ran El-Yaniv. 2017. Selective Classification for Deep Neural Networks. In *Advances in Neural Information Processing Systems*, Vol. 30. Curran Associates.
- [4] Shihao Gu, Bryan Kelly, and Dacheng Xiu. 2020. Empirical Asset Pricing via Machine Learning. *Review of Financial Studies* 33, 5 (2020), 2223–2273.
- [5] David Ha and Jürgen Schmidhuber. 2018. Recurrent World Models Facilitate Policy Evolution. In *Advances in Neural Information Processing Systems*, Vol. 31. Curran Associates.
- [6] Danijar Hafner, Timothy Lillicrap, Jimmy Ba, and Mohammad Norouzi. 2020. Dream to Control: Learning Behaviors by Latent Imagination. In *International Conference on Learning Representations*. OpenReview.net.
- [7] Campbell R. Harvey, Yan Liu, and Heqing Zhu. 2016. ... and the Cross-Section of Expected Returns. *Review of Financial Studies* 29, 1 (2016), 5–68.

- [8] Bryan T. Kelly, Seth Pruitt, and Yanan Su. 2019. Characteristics Are Covariances: A Unified Model of Risk and Return. *Journal of Financial Economics* 134, 3 (2019), 501–524.
- [9] Yann LeCun. 2022. A Path Towards Autonomous Machine Intelligence. *OpenReview preprint* (2022).
- [10] Yang Li, Yangyang Yu, Haohang Li, Zhiyang Chen, and Khaldoun Khashanah. 2023. TradingGPT: Multi-Agent System with Layered Memory and Distinct Characters for Enhanced Financial Trading Performance. *arXiv preprint arXiv:2309.03736* (2023).
- [11] Yukun Liu, Aleh Tsyvinski, and Xi Wu. 2022. Common Risk Factors in Cryptocurrency. *Journal of Finance* 77, 2 (2022), 1133–1177.
- [12] Alejandro Lopez-Lira and Yuehua Tang. 2023. Can ChatGPT Forecast Stock Price Movements? Return Predictability and Large Language Models. *arXiv preprint arXiv:2304.07619* (2023).
- [13] Igor Makarov and Antoinette Schoar. 2020. Trading and Arbitrage in Cryptocurrency Markets. *Journal of Financial Economics* 135, 2 (2020), 293–319.
- [14] Stefan Nagel. 2021. *Machine Learning in Asset Pricing*. Princeton University Press.
- [15] Dimitris N. Politis and Joseph P. Romano. 1994. The Stationary Bootstrap. *J. Amer. Statist. Assoc.* 89, 428 (1994), 1303–1313.
- [16] Halbert White. 2000. A Reality Check for Data Snooping. *Econometrica* 68, 5 (2000), 1097–1126.
- [17] Hongyang Yang, Xiao-Yang Liu, and Christina Dan Wang. 2023. FinGPT: Open-Source Financial Large Language Models. In *FinLLM Symposium at IJCAI*.
- [18] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2023. ReAct: Synergizing Reasoning and Acting in Language Models. In *International Conference on Learning Representations*. OpenReview.net.
- [19] Yangyang Yu, Haohang Li, Zhiyang Chen, Yuechen Jiang, Yang Li, Denghui Zhang, Rong Liu, Jordan W. Suchow, and Khaldoun Khashanah. 2024. FinMem: A Performance-Enhanced LLM Trading Agent with Layered Memory and Character Design. *Proceedings of the AAAI Symposium Series* 3, 1 (2024), 595–597.
- [20] Wentao Zhang, Lingxuan Zhao, Hao Xia, Shuo Sun, Jiaze Sun, Molei Qin, Xinyi Li, Yuqing Zhao, Yilei Zhao, Xinyu Cai, et al. 2024. A Multimodal Foundation Agent for Financial Trading: Tool-Augmented, Diversified, and Generalist. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*. ACM.