

The Data End of Crypto Quant for Public LLMs: An Observation Layer that Compiles World Bundles

Anonymous Author(s)

Abstract

Purpose. End-to-end crypto quant for public LLMs needs a *data end* before decision or execution. We build that layer as an anonymized observation runtime: multi-band evidence $\rightarrow$ BandPIT/WMI $\rightarrow$ a point-in-time (PIT) *world bundle*, with hard refusal when support is thin. **Evidence.** Primary—seeing: grounding recovers ready/missing/tilt at ≈ 0.97 under Compiled vs. $\lesssim 0.23$ under Raw, while sufficiency chat alone is unreliable (Compiled mean 0.60). Secondary—valve: Compiled thin-abstain is 1.0 (100-day and full OOS $n=2000$) while Ungated soft judgment is only ≈ 0.68 / ≈ 0.56 ; Raw over-trades. Secondary—handoff: scoping WMI to the declared archive opens the production gate on $\approx 61\%$ OOS days; a no-LLM tilt rule then acts (Sharpe 0.77, CE +0.25) while full-schema WMI stays capped (max ≈ 0.093). **Claim.** Build the quant data end first: compile so models can see; refuse when they must not act.

CCS Concepts

• **Computing methodologies** $\rightarrow$ **Artificial intelligence**; *Machine learning*; • **Applied computing** $\rightarrow$ *Economics*.

Keywords

data end, observation layer, world bundle, grounding, public LLMs, cryptocurrency, point-in-time, quality-gated refusal, systems

ACM Reference Format:

Anonymous Author(s). 2026. The Data End of Crypto Quant for Public LLMs: An Observation Layer that Compiles World Bundles. In *Proceedings of the 7th ACM International Conference on AI in Finance (ICAIF '26)*. ACM, New York, NY, USA, 8 pages. <https://doi.org/10.1145/XXXXXXX.XXXXXXX>

1 Introduction

Purpose (read this first). Build the *data end* of end-to-end crypto quant for public LLMs. The stack is **data/observation** $\rightarrow$ **decision** $\rightarrow$ **execution** (Table 1); agents and generative world models [6, 7, 9, 19–21] typically assume the first stage. Crypto violates that assumption: exchange, macro, and alternative evidence arrive on incompatible clocks [11, 13]. We own only the data end—compile a PIT world bundle so models can *see*, and hard-refuse when the view is too thin. Refusal is a safety valve, not the product and not a trading claim.

Failure modes. No feed $\rightarrow$ refuse (safe, useless as a data end). Thin price fragment $\rightarrow$ trade and lose (unsafe for any downstream stack). Disclosed but ungated world $\rightarrow$ read fields, disagree on action. So the data end must (i) make the market legible and (ii) enforce a quality valve before decision/execution.

Table 1: End-to-end quant stack; this paper owns only the data end.

Stage	Typical owner	This paper
Data / observation	Assumed by agents	World bundle + quality valve
Decision	LLM / policy / portfolio	Out of scope
Execution	Broker / OMS	Out of scope

What we build. An anonymized runtime (Section 4): multi-band evidence $\rightarrow$ BandPIT/WMI $\rightarrow$ JSON world bundle for GPT / DeepSeek / GLM / Gemini class models, served through a read API a decision layer can call. “World model” means state compiler + quality gate (Section 3), not dynamics and not a decision engine. The bundle is complete, honest, and auditable; `should_ai_abstain` is the machine-readable valve.

How we evaluate. **RQ1 (primary—seeing):** Are ready / missing / tilt answers verifiable against PIT ground truth under Compiled vs. Raw? **RQ2 (valve):** Does a hard flag enforce thin abstention where soft disclosure does not? Four arms: Compiled, Ungated (same content, no flag), Raw (mom5), Blind. **RQ3 (secondary):** Can a transparent no-LLM rule hand off from the bundle (production refuse vs. nonempty tilts)?

Contributions. (1) Data-end semantics: epistemic observations, compilation Π_t , world-quality index, and seeing vs. gating with a lag-missingness bound. (2) Anonymized PIT compile path: collectors $\rightarrow$ BandPIT $\rightarrow$ surface $\rightarrow$ world bundles; frozen four-arm harness. (3) Live + hand-off validation: grounding ≈ 0.97 with sufficiency dissociation; thin-abstain 1.0 under full-schema stress; scoped-archive Compiled-open handoff (open share ≈ 0.61 , Sharpe 0.77) plus ungated tilt contrast ($\Delta\text{CE}=0.334$; LOBO content share 1.0). Estimand: data-end legibility, valve reliability, and handoff usability.

2 Related Work

World models. World models learn compact states and often dynamics [6, 7, 9]. We address the complementary observation-side problem: a PIT-safe, quality-tagged crypto world *before* dynamics or LLM policy—not a Dreamer/JEPA simulator.

LLM agents in finance. Return prediction [12], open financial LLMs [18], memory- and tool-augmented agents [10, 19–21], and ICAIF retrieval/debate/judge systems [1, 3, 15, 16] improve how agents fetch, debate, and are scored. They still assume a usable observation stream at the data end of the stack. We study that prior stage: compile a world bundle, measure whether models can *see* it, then gate action when they must not act. Our estimand is therefore orthogonal to memory routers and debate judges: without a typed PIT world, those stacks consume whatever fragment the tool dump provides.

Selective prediction and PIT. Abstention can be optimal under noise [4, 5]; we attach refusal to world quality $q(W_t)$ via a runtime field, not classifier confidence alone (Prop. 3.8). Macro vintages [2] and crypto microstructure [11, 13] motivate multi-band compilation with disclosed missingness (Prop. 3.6). Bootstrap discipline for the content probe follows [8, 14, 17].

3 Observation-Layer Semantics

This section states the data-end theory as *interface semantics*: what counts as an observation, how compilation yields a world, and when quality forces refusal. It is not a pricing theorem and not a generative dynamics model. Systems realization is Section 4; empirical tests are Section 5.

Definition 3.1 (Epistemic observation). $O_{j,t} = (x_{j,t}, \tau_{j,t}, q_{j,t}, g_{j,t}, r_{j,t})$: value, latest-available time, quality, main-view gate, semantic role. Omitting (τ, q, g, r) yields a feature dump, not a world observation.

Definition 3.2 (Compilation). $\Pi_t = B_t \circ M_t \circ A_t$: align clocks, apply honesty/missingness gates, assemble band views into a consumer object.

Definition 3.3 (World bundle). $W_t = \Pi_t(\mathcal{F}_t^{\text{raw}})$ with scalar quality $q(W_t)$; $\mathcal{F}_t^{\text{AI}} = \sigma(W_t)$. W_t is the handoff object to decision.

Definition 3.4 (World-quality index). Ready/limited/missing counts disclose coverage. $\text{WMI}_t = B_t U_t H_t$ aggregates breadth, stability, and honesty in $[0, 1]$; ACWMI multiplies a regime weight. Band scope selects the support of (B, U, H) : full schema vs. the declared consumer archive. The Compiled interface exposes `should_ai_abstain` (1 if $q(W_t) < q^*$ and, under archive scope, incompleteness) with `thin_world` (production $q^* = 0.2$ for WMI).

Seeing vs. gating. The bundle is what the model sees; the hard flag *blocks* action when $q(W_t) < q^*$. Tilts without a gate are a dump; a gate without a readable bundle is only a switch. The two interfaces are separable: Ungated keeps W_t and numeric $q(W_t)$ but removes the hard boolean.

Assumption 3.5 (Bounded lag and increments). Band j has max observation lag Δ_j ; latent increments satisfy $\|S_u - S_{u-1}\| \leq \bar{\delta}$.

PROPOSITION 3.6 (LAG-MISSINGNESS RECONSTRUCTION BOUND). Under Assumption 1, any compiler that uses only in-clock observations satisfies

$$\|\tilde{S}_t - S_t\| \leq C_{\text{del}} \bar{\delta} \sum_j \Delta_j + C_{\text{noi}} \varepsilon_t + C_{\text{miss}} m_t,$$

where m_t is disclosed missing mass and ε_t collects measurement noise. Operationally, BandPIT maps age to ready/limited/missing so m_t and lag terms are readable in W_t ; compilation cannot erase delay, but it can refuse to pretend the bound is zero.

PROPOSITION 3.7 (COMPILATION $\neq$ DUMP). Enlarging raw span without Π_t need not enlarge usable $\mathcal{F}^{\text{AI}}$: ungated stale evidence can raise volume while lowering honesty, tightening rather than loosening Prop. 3.6.

PROPOSITION 3.8 (WORLD-CONDITIONAL ABSTENTION). Let $c_{\text{act}}(W_t)$ be the expected cost of non-abstain actions and c_{abs} the abstain cost. If $c_{\text{act}}(W_t) > c_{\text{abs}}$ and c_{act} is nonincreasing in $q(W_t)$, the Bayes action is abstain on a lower set $\{W : q(W) < q^*\}$ [4]. A hard runtime flag implements that set; soft text does not.

Table 2: Compile-path modules (systems view, not a product inventory).

Module	Role
Collectors / stores	Timestamped / vintaged band history
BandPIT compiler	Previous-close readiness; Π_t
Quality indices	$B, U, H \rightarrow$ WMI; abstain flag; O_t
Bundle builder	Complete / honest / auditable JSON
Service surface	Bundle / health / availability reads
Consumer harness	Four-arm protocol; temp. 0

PROPOSITION 3.9 (LOBO CONTENT VS. GATING). *Deleting band j changes a downstream value functional through a content channel and a gating channel. Under an ungated rule the gating channel is shut, so the value gap identifies nonempty fields.*

Theory $\rightarrow$ protocol. Prop. 3.6–3.7 $\Rightarrow$ RQ1 (can the consumer recover disclosed readiness/tilts?). Prop. 3.8 $\Rightarrow$ RQ2 (hard flag vs. soft judgment on the same W_t). Prop. 3.9 $\Rightarrow$ RQ3 (no-LLM handoff / LOBO content share).

REMARK 3.10 (NOT A GENERATIVE WM). *We do not learn $p(s_{t+1} | s_t, a_t)$. We deliver a state compiler with a quality-gated refuse interface—the data end of the stack.*

4 Runtime: Data-End Compile Path

The anonymized prototype implements the data end as a layered compile path (Figure 1, Table 2). Decision and execution sit downstream; this paper does not claim them. Paths and product names are omitted for double-blind review.

4.1 Layered design

- (1) **Data layer.** Exchange, macro, and alternative collectors write vintage-aware history (observation timestamps; macro `available_at`).
- (2) **Store layer.** Embedded stores hold raw series, merged market panels, and compiled analytics (readiness, WMI/ACWMI, world snapshots).
- (3) **Logic layer.** BandPIT reconstructs readiness on a previous-close clock; honesty/missingness gates implement Π_t ; WMI/ACWMI and availability shocks O_t (band outage events for quasi-exogenous levers) are first-class.
- (4) **Service layer.** A read API serves quality-tagged bundles and health objects without restarting collectors for newly arrived rows. Anonymized surface includes world-bundle and decision-bundle reads, health exposing `should_ai_abstain`, a paper world-model time slice, and availability-shock queries—the machine-readable handoff a decision layer would call.
- (5) **Consumer layer.** Frozen-prompt LLM/rule agents call an OpenAI-compatible adapter at temperature 0 under four arms.

4.2 Evaluation archive

The runtime schema is band-extensible; this study evaluates the three-band consumer object in Table 3 on a PIT panel of ≈ 3990 asset-days (2025-07 to 2026-08): exchange ready share ≈ 0.80 , macro 1.0,




Figure 1: Data-end pipeline: evidence → BandPIT → world bundle → Compiled / Ungated / Raw / Blind.

Table 3: Input evidence vs. fields emitted in the world bundle (live archive).

Band	Fetches (raw)	Emitted (bundle)
exchange	OHLCV, funding, OI	mom5; status/age
macro	equity-vol / dollar indices	macro_tilt; status/age
alternative	stablecoin circulation	alt_tilt; status/age
derived	—	WMI, ACWMI, abstain, O_t , evidence IDs

Table 4: BandPIT freshness thresholds (seconds → days).

Band	Fresh window	In eval archive
exchange	2d	yes (ready ≈0.80)
macro	5d	yes (ready 1.0)
alternative	7d	yes (ready 1.0)
news / on-chain / options	3d	missing
tokenomics	7d	missing
event calendar	14d	missing

alternative 1.0. Full-schema WMI scores permanently empty non-archive slots and caps at max ≈0.093; the production valve therefore never opens under that scoring. Scoped WMI (WORLD_MODEL_BAND_SCOPE=eval.archive) and exposes thin_world. Figure 2: macro and alternative stay ready; exchange readiness varies and defines thick/thin days. Figure 3: full-schema WMI stays under 0.2; ACWMI can look healthier under regime weight, so the runtime exports both indices and the configured band scope. Table 5: scoped production gate opens on ≈61% OOS days (≡ 3/3 archive-ready); full-schema scoring remains a closed stress baseline.

4.3 PIT clock and freshness

For calendar day t , BandPIT reconstructs the decision information set at $(t-1)$ 23:59 (payoff is same-day close-to-close return); the bundle records decision_asof. Each band is graded ready / limited / missing against daily-scale age thresholds (Table 4). Missingness is disclosed rather than imputed (Prop. 3.6). Breadth, stability, and honesty aggregate into $WMI_t = B_t U_t H_t$ (ACWMI adds regime weight). PIT-safe tilts use only in-clock observations: five-day macro

changes, a seven-day stablecoin-flow slope, and mom5. The valve sets should_ai_abstain when the scoped archive is incomplete or WMI < 0.2. Figure 2: macro and alternative stay ready; exchange readiness varies and defines thick/thin days. Figure 3: full-schema WMI stays under 0.2; ACWMI can look healthier under regime weight, so the runtime exports both indices and the configured band scope. Table 5: scoped production gate opens on ≈61% OOS days (≡ 3/3 archive-ready); full-schema scoring remains a closed stress baseline.

4.4 World bundle (what the LLM sees)

The consumer object is a JSON world bundle (Table 6): timing, completeness, honesty, quality (including the abstain flag), content tilts, and evidence_ids. Listing 1 shows an exchange-gap day




Figure 2: PIT readiness: what is visible in the archive over time.




Figure 3: World-quality paths: full-schema WMI stays under 0.2; ACWMI is regime-weighted.

Table 5: WMI gate vs. open share (OOS). Scoped archive vs. full-schema legacy.

Gate	Open share	Note
0.2 scoped archive	0.61	complete $\wedge$ WMI $\geq$ 0.2
0.2 full-schema legacy	0.00	max WMI $\approx$ 0.093
0.05 WMI-only (scoped)	1.00	no completeness $\wedge$

Table 6: World-bundle field groups (the market view).

Group	Fields
Timing	decision_asof, previous-close protocol
Completeness	ready/limited/missing counts, band status
Honesty	B, U, H ; main-view / stale gates
Quality	WMI, ACWMI, abstain flag, thin_world
Content	macro_tilt, alt_tilt, mom5
Audit	evidence_ids; EAR required

(scoped valve closed); Listing 2 shows a 3/3-ready day where scoped WMI clears 0.2 and should_ai_abstain=false, while full-schema WMI remains disclosed as a ceiling audit.

Listing 1: Exchange-gap bundle (Compiled; scoped closed).

```
{
  "date": "2026-05-11", "asset": "ADA",
  "decision_asof": "2026-05-10T23:59:00",
  "completeness": {"n_ready": 2, "n_missing": 1},
  "world_model_index": {"band_scope": "eval_archive",
    "wmi": 0.139, "archive_complete": false,
    "should_ai_abstain": true, "thin_world": true,
    "full_schema_wmi": 0.015},
  "macro_tilt": 1.0, "alt_tilt": -1.0}
```

Listing 2: 3/3-ready day: scoped valve opens.

```
{
  "date": "2026-01-16", "asset": "ADA",
  "completeness": {"n_ready": 3, "n_missing": 0},
  "world_model_index": {"band_scope": "eval_archive",
    "wmi": 1.0, "archive_complete": true,
    "should_ai_abstain": false, "thin_world": false,
    "full_schema_wmi": 0.093},
  "band_status": {"exchange": "ready", "macro": "ready",
    "alternative": "ready"},
  "macro_tilt": 1.0, "alt_tilt": -1.0}
```

4.5 Four arms and frozen prompts

Actions $\in$ {bullish, bearish, neutral, abstain}. **Compiled**: full bundle; frozen prompt *requires* abstain when should_ai_abstain is true (and when the world is judged thin). **Ungated**: same content, completeness, and numeric WMI; *no* hard flag; soft prompt asks the model to judge thinness. **Raw**: mom5 only—thin price wiring for the ladder. **Blind**: “Today is d . How should I trade a ?” with no feed. Live LLM arms in Section 5.3 were frozen under full-schema scoring (max WMI $\approx$ 0.093), so that Compiled path abstains on every OOS day—a closed-valve stress test. The scoped-archive path (default band scope) opens on band-thick days and is the handoff contract in Section 5.4. Compiled vs. Ungated differ only in enforcement (Listings 3–4); prompts, IS/OOS cut, and seeds are frozen. On thin Compiled days the evidence-alignment proxy (EAR) is 1.0 across vendors: rationales stay bound to disclosed evidence_ids.

Listing 3: Compiled (excerpt): hard valve.

1. If should_ai_abstain is true OR the world is thin, you MUST "abstain".
2. Prefer band content over raw momentum (macro_tilt, alt_tilt are PIT-safe).
3. Rationale must match evidence_ids.

Listing 4: Ungated (excerpt): soft judgment.

1. Judge whether low WMI / missingness makes action unsafe; you MAY "abstain".
2. Prefer band content when you act.
- NO hard abstain boolean; NO forced policy.

Why the hard valve is not a confound. Compiled abstention is partly prompt obedience, and that is intentional: the data end’s job on thin days is to *block* action, not to hope the model discovers thinness (Prop. 3.8). Ungated measures soft judgment from the same disclosed bundle; Raw measures fragment wiring; Blind measures no feed. The gating result is that the hard flag is reliable across vendors; soft refusal is not—which is why half-baked context remains dangerous without the valve.

4.6 Systems validation (engines)

Paper-lab smoke against the production engines returns a wired WMI/ACWMI sample with should_ai_abstain consistent with the configured index mode (WMI gate 0.2; ACWMI gate 0.35). The

Table 7: Grounding (50 days): primary seeing result. C = Compiled; R = Raw.

Model	Ready F1		Missing F1		Tilt-sign	
	C	R	C	R	C	R
gpt-5.4-mini	0.94	0.00	0.94	0.31	0.94	0.20
deepseek-v4-flash	0.98	0.06	0.98	0.20	0.98	0.18
glm-5.2	0.98	0.09	1.00	0.11	0.98	0.20
gemini-3.5-flash-lite	0.96	0.03	0.96	0.30	0.96	0.18
Mean	0.97	0.04	0.97	0.23	0.97	0.19

same indices appear in the consumer bundle; availability shocks O_t are queryable on the service surface.

5 Experiments

5.1 Setup

PIT panel ~ 399 days $\times 10$ liquid names (≈ 3990 asset-days; 2025-07–2026-08), chronological IS/OOS 200/200 days (cut 2026-01-16). Unless noted, live arms share the same stratified 100 OOS asset-days (seed fixed), temperature 0, and four flash/mini vendors (gpt-5.4-mini, deepseek-v4-flash, glm-5.2, gemini-3.5-flash-lite) via an OpenAI-compatible adapter. Grounding uses 50 OOS days; full-OOS abstention uses $n=2000$; the no-LLM probe uses 199 OOS days after return alignment. Primary metrics: RQ1 grounding F1/tilt accuracy (sufficiency chat secondary); RQ2 thin-world abstain rate (CE as avoided loss, not alpha); RQ3 Δ CE of the content rule vs. momentum.

5.2 RQ1 (primary): Seeing—grounding workflow

If the data end works, models must report a *checkable* market view—not invent coverage from a price fragment.

On 50 OOS asset-days, each model lists ready/missing bands among exchange/macro/alternative, reports macro_tilt / alt_tilt signs, and judges sufficiency; (i)–(iii) score against PIT ground truth.

Table 7: Compiled mean ready/missing F1 and tilt-sign accuracy 0.97; Raw collapses to 0.04, 0.23, 0.19. Raw has no readiness or tilt fields, so it cannot recover archive state; Compiled forces use of the disclosed world.

Legibility $\neq$ sufficiency chat. Table 8: Compiled sufficiency accuracy averages only 0.60 (Gemini 0.04) despite 0.97 field grounding, while Raw often answers “enough?” correctly by chance/heuristics (0.95) while inventing coverage. Chat judgment of “is this enough to trade?” is not a substitute for a typed world; we therefore measure seeing by checkable fields (RQ1) and action safety by the hard valve (RQ2).

Case study. Listings 5–6: ADA (2026-05-11) exchange missing; Compiled recovers the checkable view; Raw cannot. A second day (2026-01-17) shows the valve contrast: Compiled abstains on should_ai_abstain; Raw goes bullish from mom5 = +1 alone.

Listing 5: Compiled grounding (gpt-5.4-mini, ADA 2026-05-11).


```
ready: alternative, macro
```

Table 8: Sufficiency chat vs. field grounding (50 days). Seeing = fields; chat alone misleads.

Model	Suff. C	Suff. R	Ready F1 C	Ready F1 R
gpt-5.4-mini	0.78	0.98	0.94	0.00
deepseek-v4-flash	0.78	0.90	0.98	0.06
glm-5.2	0.80	0.96	0.98	0.09
gemini-3.5-flash-lite	0.04	0.94	0.96	0.03
Mean	0.60	0.95	0.97	0.04

Table 9: Thin-world abstention on identical 100 OOS days (Wilson 95% CIs). Blind = 1.0.

Model	Thin C	Thin U [CI]	Abs. R [CI]	Δ CE
gpt-5.4-mini	1.00	0.63 [.53,.72]	0.04 [.02,.10]	+0.73
deepseek-v4-flash	1.00	0.81 [.72,.87]	0.75 [.66,.82]	+1.22
glm-5.2	1.00	0.86 [.78,.91]	0.11 [.06,.19]	+1.04
gemini-3.5-flash-lite	1.00	0.43 [.34,.53]	0.07 [.03,.14]	+1.30
Mean	1.00	0.68	—	+1.07

**Figure 4: Safety valve: Blind refuses; Raw over-trades; Ungated is unstable; Compiled enforces the hard gate.**

```
missing: exchange
macro_tilt +; alt_tilt -; sufficiency: insufficient
"thin world... should_ai_abstain=true"
```

Listing 6: Raw grounding / action (same model family).

```
grounding: ready=[]; cannot confirm tilts
action (2026-01-17): bullish
"Positive 5-period momentum in raw feed."
```

5.3 RQ2 (safety valve): Hard refuse when the view is thin

Legibility does not imply safe action. Table 9 and Figure 4: on the same 100 days, Compiled abstains on every thin day (1.0); Ungated soft-judges at mean 0.68 (Gemini 0.43); Raw over-trades (Table 10); Blind abstains 1.0 with no feed—safe but useless as a data end. The costly failure mode is the middle case: half-baked Raw context makes models act and lose. Mean Δ CE (Compiled–Raw) +1.07 is avoided loss under the valve, not an alpha claim.

Table 10: Raw action mix (100 days): half-baked context drives directional mass.

Model	Abs.	Bull.	Bear.	Neut.
gpt-5.4-mini	0.04	0.28	0.38	0.30
deepseek-v4-flash	0.75	0.14	0.04	0.07
glm-5.2	0.11	0.40	0.48	0.01
gemini-3.5-flash-lite	0.07	0.44	0.49	0.00

Table 11: Full-OOS thin abstention ($n=2000$ asset-days; all thin under gate 0.2).

Model	Thin C	Thin U	Abs. R
gpt-5.4-mini	1.000	0.666	0.000
deepseek-v4-flash	1.000	0.540	0.748
glm-5.2	1.000	0.679*	0.078*
gemini-3.5-flash-lite	1.000	0.359	0.000
Mean	1.000	≈ 0.56	—

*GLM coverage incomplete (Ungated $n=1402$; Raw $n=675$).

Table 12: Open slice (3/3 ready $\equiv$ WMI ≥ 0.05 ; $n \approx 1224$). Field lower bound under soft judgment.

Model	Abs. U	CE U	Act-CE U	Abs. R	CE R
gpt-5.4-mini	0.58	+0.42	+0.32	0.00	-0.28
deepseek-v4-flash	0.42	+0.21	+0.45	0.75	+0.04
glm-5.2	0.58	+0.24	+0.11	0.09*	-0.87*
gemini-3.5-flash-lite	0.34	+0.31	+0.06	0.00	-0.25

*Raw GLM $n=398$; Ungated GLM $n=932$; others $n=1224$. Bootstrap mean CE $\approx +0.31$ ($B=5,999$ reps; one-sided $p \approx 0.21$).

Full-OOS replication. Table 11: on all 2000 OOS asset-days (every day thin under the 0.2 gate), Compiled thin-abstain remains 1.000 (Wilson [.998, 1]) for every vendor. Ungated soft rates fall to a four-vendor mean ≈ 0.56 ; Raw GPT/Gemini abstain at 0.000. Scale does not rescue soft judgment. Blind refuses everywhere with no feed on the paired 100-day sample (abstain 1.0 for every vendor): the failure mode that needs a data end is not “models never abstain,” but “thin fragments make them act.”

Inside Ungated soft acts. On the mixed 100-day sample, when Ungated *does* act, act-conditional CE is sharply negative for three of four vendors (DeepSeek -3.12 , GLM -4.28 , Gemini -1.67) while GPT is positive on its smaller act set. Soft disclosure without a hard flag is not a free lunch: the same readable world that RQ1 scores can still harm decisions when the valve is removed.


Open slice. Table 12: on WMI ≥ 0.05 ($\approx 61\%$ OOS; $n \approx 1224$) Ungated keeps compiled content without the hard boolean—mean abs. ≈ 0.48 , CE positive for every vendor (mean $\approx +0.29$; bootstrap $\approx +0.31$, CIs include 0). Production Compiled still abstains at 1.0; Raw stays risky (GPT/Gemini abs. 0.0, CE negative). Point-positive open-slice CE and negative mixed-sample act-conditional CE can coexist: regime and act selection matter.

Table 13: Handoff backtest (no LLM). Scoped Compiled-open vs. full-schema closed.

Policy (reads bundle fields)	Sharpe	CE	Abstain
Full-schema gate 0.2	0.000	0.000	1.00
Scoped gate 0.2 + tilts	0.772	+0.253	0.39
Momentum (mom5)	0.101	-0.075	0.00
Buy-and-hold	-1.399	-0.987	0.00

Table 14: No-LLM content contrast. Bootstrap Δ CE 0.334; raw gap 0.336.

Policy	Sharpe	CE
Buy-and-hold	-1.404	-1.163
Momentum	0.096	-0.206
World-bundle tilts	0.764	+0.130
Tilts – Momentum	—	0.334 ($p=0.034$)

**Figure 5: Secondary content probe: compiled tilts beat momentum; Δ CE positive for every asset.**

5.4 RQ3 (secondary): Handoff backtest—Compiled-open proof

The data end must be usable by a downstream decision rule without an LLM in the loop. Full-schema WMI never clears 0.2 on this panel, so a legacy production gate alone cannot demonstrate handoff. Scoped archive WMI opens on $\approx 61\%$ OOS days ($\equiv$ 3/3 ready); DecisionHandoffService then abstains when the valve is closed and otherwise acts on disclosed macro_tilt / alt_tilt. Table 13: scoped Compiled-open handoff Sharpe 0.772 / CE +0.253 (abstain 0.39) vs. momentum 0.101 / -0.075 and buy-and-hold -1.40; full-schema legacy remains fully closed. Table 14 and Figure 5 keep the ungated content contrast (tilts vs. momentum; Δ CE=0.334, $p=0.034$; LOBO content share 1.0) as a nonempty-field probe on the same vintaged tilts.

Density split and LOBO channel. Table 15: open/thick Δ CE=0.315; exchange-gap days abstain more (Δ CE=0.129). Table 16 and Figure 6


Table 15: Content probe by archive density (OOS).

Slice	Share	Mech CE	Δ CE vs mom
All OOS	1.00	0.130	0.334*
Open / thick	0.61	0.061	0.315
Thin (exchange gap)	0.39	1.048	0.129

*Block-bootstrap Δ CE; raw mean gap 0.336.

Table 16: LOBO decomposition (ungated content probe). Content share 1.0.

Drop	Δ CE content	p	Content share
macro	-0.334	0.034	1.0
alternative	-0.334	0.034	1.0
exchange	0.000	1.000	—

**Figure 6: LOBO: deleting macro/alternative removes the edge via content, not the refuse flag.**

(Prop. 3.9): deleting macro or alternative collapses the gap through the content channel (share 1.0; gating residual 0).

6 Discussion

What carries the paper. (i) The data end is legible: grounding ≈ 0.97 , with sufficiency chat a poor substitute (RQ1 / Props. 3.6–3.7). (ii) Thin views must not drive action: soft text fails at 100-day and full-OOS scale; Blind refuses with no feed; Raw acts on fragments and loses; the hard valve works (RQ2 / Prop. 3.8). (iii) Scoped-archive WMI opens the production gate on band-thick days; a no-LLM handoff then acts on disclosed tilts while full-schema scoring stays closed (RQ3 / Prop. 3.9); LOBO content share 1.0. Relative to FinMem /FinAgent and ICAIF agent stacks [1, 3, 15, 16, 20, 21], we supply the observation stage those systems assume.

Architecture. Semantics (Section 3) define the interfaces; Runtime (Section 4) implements Π_r and the service handoff; Experiments (Section 5) test seeing, gating, and nonempty fields in that order. Trading PnL is intentionally secondary: it confounds observation quality with policy skill.

Design. World bundle not feature dump; previous-close PIT with explicit freshness; four arms separate seeing+gating, seeing without gating, thin wiring, and no feed. Service-layer health and bundle

reads make the valve and the view callable by any downstream policy.

Reproducibility. Paper-lab smoke checks BandPIT/WMI/ACWMI wiring; PIT rebuild and the frozen four-arm harness regenerate the panel and live transcripts used here. Prompts, IS/OOS cut, seeds, and vendor IDs are fixed; credentials stay in environment variables. An anonymized harness ships on acceptance.

Scope. We evaluate the data end on the PIT archive above; decision and execution sit downstream.

7 Conclusion

End-to-end crypto quant for public LLMs needs a reliable data end before decision or execution. We contribute that observation layer: compile asynchronous evidence into a consumable PIT world bundle, disclose structural missingness, and hard-refuse when the view is too thin. Grounding recovers ready/missing/tilt at ≈ 0.97 under Compiled vs. Raw collapse; sufficiency chat alone fails; the hard valve abstains 1.0 on thin days where soft judgment does not; scoping WMI to the declared archive opens Compiled handoff on band-thick days (Sharpe 0.77) while full-schema WMI remains a closed ceiling. Build the quant data end first: compile so models can see; refuse when they must not act.

References

- [1] Chanyeol Choi, Jihoon Kwon, Alejandro Lopez-Lira, Chaewon Kim, Minjae Kim, Junha Hwang, Jaeseon Ha, Hojun Choi, Suyeol Yun, Yongjin Kim, and Yongjae Lee. 2025. FinAgentBench: A Benchmark Dataset for Agentic Retrieval in Financial Question Answering. In *Proceedings of the 6th ACM International Conference on AI in Finance*. ACM, 632–637. doi:10.1145/3768292.3770362
- [2] Dean Croushore and Tom Stark. 2001. A Real-Time Data Set for Macroeconomists. *Journal of Econometrics* 105, 1 (2001), 111–130.
- [3] Yitong Duan, Chuheng Zhang, and Jian Li. 2025. FactorMAD: A Multi-Agent Debate Framework Based on Large Language Models for Interpretable Stock Alpha Factor Mining. In *Proceedings of the 6th ACM International Conference on AI in Finance*. ACM, 605–613.
- [4] Ran El-Yaniv and Yair Wiener. 2010. On the Foundations of Noise-Free Selective Classification. *Journal of Machine Learning Research* 11 (2010), 1605–1641.
- [5] Yonatan Geifman and Ran El-Yaniv. 2017. Selective Classification for Deep Neural Networks. In *Advances in Neural Information Processing Systems*, Vol. 30. Curran Associates.
- [6] David Ha and Jürgen Schmidhuber. 2018. Recurrent World Models Facilitate Policy Evolution. In *Advances in Neural Information Processing Systems*, Vol. 31. Curran Associates.
- [7] Danijar Hafner, Timothy Lillicrap, Jimmy Ba, and Mohammad Norouzi. 2020. Dream to Control: Learning Behaviors by Latent Imagination. In *International Conference on Learning Representations*. OpenReview.net.
- [8] Campbell R. Harvey, Yan Liu, and Heqing Zhu. 2016. ... and the Cross-Section of Expected Returns. *Review of Financial Studies* 29, 1 (2016), 5–68.
- [9] Yann LeCun. 2022. A Path Towards Autonomous Machine Intelligence. *OpenReview preprint* (2022).
- [10] Yang Li, Yangyang Yu, Haohang Li, Zhiyang Chen, and Khaldoun Khashanah. 2023. TradingGPT: Multi-Agent System with Layered Memory and Distinct Characters for Enhanced Financial Trading Performance. *arXiv preprint arXiv:2309.03736* (2023).
- [11] Yukun Liu, Aleh Tsyvinski, and Xi Wu. 2022. Common Risk Factors in Cryptocurrency. *Journal of Finance* 77, 2 (2022), 1133–1177.
- [12] Alejandro Lopez-Lira and Yuehua Tang. 2023. Can ChatGPT Forecast Stock Price Movements? Return Predictability and Large Language Models. *arXiv preprint arXiv:2304.07619* (2023).
- [13] Igor Makarov and Antoinette Schoar. 2020. Trading and Arbitrage in Cryptocurrency Markets. *Journal of Financial Economics* 135, 2 (2020), 293–319.
- [14] Dimitris N. Politis and Joseph P. Romano. 1994. The Stationary Bootstrap. *J. Amer. Statist. Assoc.* 89, 428 (1994), 1303–1313.
- [15] Yiqing Shen, Jingshu Zhang, Feng Chen, Kaiyuan Yan, and Hongguang Li. 2025. FinSearch: A Temporal-Aware Search Agent Framework for Real-Time Financial Information Retrieval with Large Language Models. In *Proceedings of the 6th ACM International Conference on AI in Finance*. ACM. doi:10.1145/3768292.3770382

- [16] Rui Sun, Zuo Bai, Wentao Zhang, Yuxiang Zhang, Li Zhao, Shan Sun, and Zhengwen Qiu. 2025. FinResearchBench: A Logic Tree based Agent-as-a-Judge Evaluation Framework for Financial Research Agents. In *Proceedings of the 6th ACM International Conference on AI in Finance*. ACM, 656–664. doi:10.1145/3768292.3770364
- [17] Halbert White. 2000. A Reality Check for Data Snooping. *Econometrica* 68, 5 (2000), 1097–1126.
- [18] Hongyang Yang, Xiao-Yang Liu, and Christina Dan Wang. 2023. FinGPT: Open-Source Financial Large Language Models. In *FinLLM Symposium at IJCAI*.
- [19] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2023. ReAct: Synergizing Reasoning and Acting in Language Models. In *International Conference on Learning Representations*. OpenReview.net.
- [20] Yangyang Yu, Haohang Li, Zhiyang Chen, Yuechen Jiang, Yang Li, Denghui Zhang, Rong Liu, Jordan W. Suchow, and Khaldoun Khashanah. 2024. FinMem: A Performance-Enhanced LLM Trading Agent with Layered Memory and Character Design. *Proceedings of the AAAI Symposium Series* 3, 1 (2024), 595–597.
- [21] Wentao Zhang, Lingxuan Zhao, Hao Xia, Shuo Sun, Jiaze Sun, Molei Qin, Xinyi Li, Yuqing Zhao, Yilei Zhao, Xinyu Cai, et al. 2024. A Multimodal Foundation Agent for Financial Trading: Tool-Augmented, Diversified, and Generalist. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*. ACM.