

See First, Refuse When Thin: Compiling Point-in-Time World Bundles for LLM Agents in Crypto Markets

Anonymous Author(s)

Abstract

Purpose. In AI-based finance, the performance frontier is often set less by model complexity than by whether the market world fed to the model is broad, stable, and honest. We treat the missing piece as a *data end*: an anonymized observation runtime that compiles multi-band crypto evidence into a point-in-time (PIT) *world bundle* (BandPIT/WMI) with hard refusal when support is thin—infrastructure for LLM/agent market analysis and for evaluating whether models *understand* that world, not a trading engine. **Evidence.** Primary—seeing: grounding recovers ready/missing/tilt at ≈ 0.97 under Compiled vs. ≤ 0.23 under Raw, while sufficiency chat alone is unreliable (Compiled mean 0.60). Secondary—valve: under a *frozen full-schema* stress gate (max WMI $\approx 0.093 < 0.2$), Compiled thin-abstain is 1.0 (100-day and full OOS $n=2000$) while Ungated soft judgment is only $\approx 0.68 / \approx 0.56$; Raw over-trades. Secondary—handoff: the *scoped-archive production valve* opens on $\approx 61\%$ OOS days; live LLMs abstain 1.0 on closed days and act on $\approx 89\%$ of open days, while a no-LLM tilt rule remains nonempty (Sharpe 0.77). **Claim.** Move from model-first to data-end-first: compile so models can see; refuse when they must not act.

CCS Concepts

• **Computing methodologies** → **Artificial intelligence**; *Machine learning*; • **Applied computing** → *Economics*.

Keywords

AI in finance, LLM agents, data end, observation layer, world bundle, grounding, market understanding, cryptocurrency, point-in-time, quality-gated refusal

ACM Reference Format:

Anonymous Author(s). 2026. See First, Refuse When Thin: Compiling Point-in-Time World Bundles for LLM Agents in Crypto Markets. In *Proceedings of the 7th ACM International Conference on AI in Finance (ICAIF '26)*. ACM, New York, NY, USA, 8 pages. <https://doi.org/10.1145/XXXXXXX.XXXXXXX>

1 Introduction

Purpose (read this first). At the methodology–application interface of AI and finance, many LLM/agent failures are not “the model is not smart enough,” but “the market worldview is too thin.” The stack is **data/observation** → **decision** → **execution** (Table 1); agents and generative world models [6, 7, 9, 19–21] typically assume the first stage. Crypto violates that assumption: exchange, macro, and alternative evidence arrive on incompatible clocks [11, 13]. We build the missing substrate—compile a PIT world bundle so models

Table 1: AI-finance stack; this paper supplies the data end for LLM/agent consumers.

Stage	Typical owner	This paper
Data / observation	Assumed by agents	World bundle + quality valve
Decision	LLM / policy / portfolio	Out of scope
Execution	Broker / OMS	Out of scope

can *see* the market, evaluate that seeing with checkable grounding, and hard-refuse when the view is too thin. The product is reduced world-model error and auditable AI judgments, not trading answers.

Failure modes. No feed → refuse (safe, useless as a data end). Thin price fragment → trade and lose (unsafe for any downstream agent). Disclosed but ungated world → read fields, disagree on action. So the data end must (i) make the market legible to LLMs/agents and (ii) enforce a quality valve before decision/execution.

What we build. An anonymized *data-provision runtime for public LLMs and finance agents* (Section 4): multi-band evidence → BandPIT/WMI → JSON world bundle for GPT /DeepSeek /GLM /Gemini-class consumers, served through a read API. It is also an evaluation substrate: the same bundle supports grounding and gated-action protocols that test whether a model understands disclosed market state. “World model” means state compiler + quality gate (Section 3), not dynamics and not a decision engine. The bundle is complete, honest, and auditable; `should_ai_abstain` is the machine-readable valve.

How we evaluate (dual protocol). **RQ1 (primary—seeing):** Are ready /missing /tilt answers verifiable against PIT ground truth under Compiled vs. Raw? **RQ2 (valve):** Under a *frozen full-schema stress gate* (WMI over all schema bands; never clears 0.2 on this panel), does a hard flag enforce thin abstention where soft disclosure does not? Four arms: Compiled, Ungated (same content, no flag), Raw (mom5), Blind. **RQ3a (handoff):** Under the *scoped-archive production valve*, do live LLMs obey closed days and act when open? Does a transparent no-LLM rule remain callable? **RQ3b (content):** On the same vintaged tilts, does an ungated content probe beat momentum (ΔCE ; LOBO)?

Contributions. (1) AI–finance data-end semantics: epistemic observations, compilation Π_t , world-quality index with band scope, and seeing vs. gating—interfaces that let LLM/agent stacks consume crypto markets safely. (2) An anonymized data-provision path for agents: collectors→BandPIT→service surface (bundle /health /handoff)→world bundles, plus a frozen four-arm harness to *evaluate* market understanding. (3) Live LLM evidence that the substrate works: grounding ≈ 0.97 ; stress-gate abstain 1.0; scoped production valve obedience

(closed 1.0; open act ≈ 0.89); nonempty tilt handoff (Sharpe 0.77; LOBO 1.0). Estimand: whether public models can see a typed market world, refuse when thin, and hand off when ready—not portfolio alpha.

2 Related Work

World models. World models learn compact states and often dynamics [6, 7, 9]. We address the complementary observation-side problem: a PIT-safe, quality-tagged crypto world *before* dynamics or LLM policy—not a Dreamer/JEPA simulator.

LLM agents in finance (ICAIF intersection). Return prediction [12], open financial LLMs [18], memory- and tool-augmented agents [10, 19–21], and ICAIF retrieval/debate/judge systems [1, 3, 15, 16] improve how agents fetch, debate, and are scored. Methodologically they still assume a usable observation stream; application stacks then conflate bad market context with bad policy. We target that gap in a data-centric sense: compile a typed PIT world for LLM/agent consumers, and score market understanding (grounding, evidence-bound refusal) before debating alpha. The estimand is orthogonal to memory routers and debate judges: without a typed world, those systems evaluate agents on tool fragments, not on whether the market view was broad, stable, and honest.

Selective prediction and PIT. Abstention can be optimal under noise [4, 5]; we attach refusal to *world* quality $q(W_t)$ via a runtime field, not classifier confidence alone (Prop. 3.8)—blocking the overconfident-when-thin failure mode. Macro vintages [2] and crypto microstructure [11, 13] motivate multi-band compilation with disclosed missingness (Prop. 3.6). Bootstrap discipline for the content probe follows [8, 14, 17].

3 Observation-Layer Semantics

This section states the data-end theory as *interface semantics*: what counts as an observation, how compilation yields a world, and when quality forces refusal. It is not a pricing theorem and not a generative dynamics model. Systems realization is Section 4; empirical tests are Section 5.

Definition 3.1 (Epistemic observation). $O_{j,t} = (x_{j,t}, \tau_{j,t}, q_{j,t}, g_{j,t}, r_{j,t})$: value, latest-available time, quality, main-view gate, semantic role. Omitting (τ, q, g, r) yields a feature dump, not a world observation.

Definition 3.2 (Compilation). $\Pi_t = B_t \circ M_t \circ A_t$: align clocks, apply honesty/missingness gates, assemble band views into a consumer object.

Definition 3.3 (World bundle). $W_t = \Pi_t(\mathcal{F}_t^{\text{raw}})$ with scalar quality $q(W_t)$; $\mathcal{F}_t^{\text{AI}} = \sigma(W_t)$. W_t is the handoff object to decision.

Definition 3.4 (World-quality index). Ready/limited/missing counts disclose coverage. $\text{WMI}_t = B_t U_t H_t$ aggregates breadth, stability, and honesty in $[0, 1]$; ACWMI multiplies a regime weight. Band scope selects the support of (B, U, H) : *full schema* (stress / ceiling audit) vs. the *declared consumer archive* (production valve). Scoped WMI scores only declared bands; it is a contract-completeness index, not a lowered q^* . The Compiled interface exposes $\text{should_ai_abstain} = 1\{q(\mathcal{V}) < q^*\}$ (and, under archive scope, incompleteness) with thin_world (production $q^* = 0.2$ for scoped WMI).

Seeing vs. gating. The bundle is what the model *sees*; the hard flag *blocks* action when $q(W_t) < q^*$. Tilts without a gate are a dump; a gate without a readable bundle is only a switch. The two interfaces are separable: Ungated keeps W_t and numeric $q(W_t)$ but removes the hard boolean.

Assumption 3.5 (Bounded lag and increments). Band j has max observation lag Δ_j ; latent increments satisfy $\|S_u - S_{u-1}\| \leq \tilde{\delta}$.

PROPOSITION 3.6 (LAG-MISSINGNESS RECONSTRUCTION BOUND). *Under Assumption 1, any compiler that uses only in-clock observations satisfies*

$$\|\tilde{S}_t - S_t\| \leq C_{\text{del}} \tilde{\delta} \sum_j \Delta_j + C_{\text{noi}} \varepsilon_t + C_{\text{miss}} m_t,$$

where m_t is disclosed missing mass and ε_t collects measurement noise. Operationally, *BandPIT* maps age to ready/limited/missing so m_t and lag terms are readable in W_t ; compilation cannot erase delay, but it can refuse to pretend the bound is zero.

PROPOSITION 3.7 (COMPILATION $\neq$ DUMP). *Enlarging raw span without Π_t need not enlarge usable $\mathcal{F}^{\text{AI}}$: ungated stale evidence can raise volume while lowering honesty, tightening rather than loosening Prop. 3.6.*

PROPOSITION 3.8 (WORLD-CONDITIONAL ABSTENTION). *Let $c_{\text{act}}(W_t)$ be the expected cost of non-abstain actions and c_{abs} the abstain cost. If $c_{\text{act}}(W_t) > c_{\text{abs}}$ and c_{act} is nonincreasing in $q(W_t)$, the Bayes action is abstain on a lower set $\{W : q(W) < q^*\}$ [4]. A hard runtime flag implements that set; soft text does not.*

PROPOSITION 3.9 (LOBO CONTENT VS. GATING). *Deleting band j changes a downstream value functional through a content channel and a gating channel. Under an ungated rule the gating channel is shut, so the value gap identifies nonempty fields.*

Theory $\rightarrow$ protocol. Prop. 3.6–3.7 $\Rightarrow$ RQ1 (recover disclosed readiness/tilts). Prop. 3.8 $\Rightarrow$ RQ2 (hard flag vs. soft judgment under the full-schema stress gate). Def. 3.4 (band scope) + Prop. 3.8 $\Rightarrow$ RQ3a (scoped production valve opens; handoff acts). Prop. 3.9 $\Rightarrow$ RQ3b (ungated content / LOBO share).

REMARK 3.10 (NOT A GENERATIVE WM). *We do not learn $p(s_{t+1} | s_t, a_t)$. We deliver a state compiler with a quality-gated refuse interface—the data end of the stack.*

4 Runtime: Data-End Compile Path

The anonymized prototype implements the data end as a layered compile path for LLM/agent consumers (Figure 1, Table 2). Decision and execution sit downstream; this paper supplies their market context and the protocols that test whether that context was understood. Paths and product names are omitted for double-blind review.

4.1 Layered design

- (1) **Data layer.** Exchange, macro, and alternative collectors write vintage-aware history (observation timestamps; macro `available_at`).
- Store layer.** Embedded stores hold raw series, merged market panels, and compiled analytics (readiness, WMI/ACWMI, world snapshots).

Table 2: Compile-path modules (systems view, not a product inventory).

Module	Role
Collectors / stores	Timestamped / vintaged band history
BandPIT compiler	Previous-close readiness; Π_t
Quality indices	$B, U, H \rightarrow$ WMI; band scope; abstain; O_t
Bundle builder	Complete / honest / auditable JSON
Service surface	Bundle / health / handoff / availability reads
Consumer harness	Four-arm protocol; temp. 0

Table 3: Input evidence vs. fields emitted in the world bundle (live archive).

Band	Fetches (raw)	Emitted (bundle)
exchange	OHLCV, funding, OI	mom5; status/age
macro	equity-vol / dollar indices	macro_tilt; status/age
alternative	stablecoin circulation	alt_tilt; status/age
derived	—	WMI, ACWMI, abstain, O_t , evidence IDs

- (3) **Logic layer.** BandPIT reconstructs readiness on a previous-close clock; honesty/missingness gates implement Π_t ; WMI/ACWMI and availability shocks O_t (band outage events for quasi-exogenous levers) are first-class.
- (4) **Service layer.** A read API serves quality-tagged bundles and health objects without restarting collectors for newly arrived rows. Anonymized surface includes world-bundle reads, `/bundle/{symbol}/handoff` (valve-gated tilt action, no LLM), health exposing `should_ai_abstain` with `band_scope`, a paper world-model time slice, and availability-shock queries.
- (5) **Consumer / eval layer.** Frozen-prompt public LLMs (and rule agents) call an OpenAI-compatible adapter at temperature 0 under four arms—the harness that scores market understanding against the supplied world.

4.2 Evaluation archive

The runtime schema is band-extensible; this study evaluates the three-band consumer object in Table 3 on a PIT panel of ≈ 3990 asset-days (2025-07 to 2026-08): exchange ready share ≈ 0.80 panel-wide (≈ 0.61 OOS), macro 1.0, alternative 1.0. **Full-schema stress WMI** scores permanently empty non-archive slots and caps at $\max \approx 0.093$, so a 0.2 gate never opens—the frozen LLM-arm protocol (RQ2). **Scoped-archive production WMI** renormalizes B, U, H to $\{\text{exchange, macro, alternative}\}$ and opens iff the archive is complete and $\text{WMI} \geq 0.2$ (RQ3a); this measures contract readiness, not a relaxed threshold.

4.3 PIT clock and freshness

For calendar day t , BandPIT reconstructs the decision information set at $(t-1)$ 23:59 (payoff is same-day close-to-close return); the bundle records `decision_asof`. Each band is graded ready / limited / missing against daily-scale age thresholds (Table 4). Missingness is disclosed rather than imputed (Prop. 3.6). Breadth, stability, and honesty aggregate into $\text{WMI}_t = B_t U_t H_t$ (ACWMI adds regime

Table 4: BandPIT freshness thresholds (seconds $\rightarrow$ days).

Band	Fresh window	In eval archive
exchange	2d	yes (ready ≈ 0.80)
macro	5d	yes (ready 1.0)
alternative	7d	yes (ready 1.0)
news / on-chain / options	3d	missing
tokenomics	7d	missing
event calendar	14d	missing

Table 5: Dual gates vs. open share (OOS). Production = scoped; stress = full schema.

Gate	Open share	Note
Scoped production 0.2	0.61	complete $\wedge$ WMI ≥ 0.2
Full-schema stress 0.2	0.00	max WMI ≈ 0.093

Table 6: World-bundle field groups (the market view).

Group	Fields
Timing	<code>decision_asof</code> , previous-close protocol
Completeness	ready/limited/missing counts, band status
Honesty	B, U, H ; main-view / stale gates
Quality	WMI, ACWMI, <code>band_scope</code> , <code>abstain</code> , <code>thin_world</code>
Content	<code>macro_tilt</code> , <code>alt_tilt</code> , <code>mom5</code>
Audit	<code>evidence_ids</code> ; <code>EAR</code> ; <code>full_schema_wmi</code>


weight). PIT-safe tilts use only in-clock observations: five-day macro changes, a seven-day stablecoin-flow slope, and `mom5`. The *scoped production valve* sets `should_ai_abstain` when the archive is incomplete or scoped WMI < 0.2 , and exposes `thin_world`; bundles also disclose `full_schema_wmi` as a ceiling audit. Figure 2: macro and alternative stay ready; exchange readiness varies and defines thick/thin days. Figure 3: same panel, two scopes—full-schema stress WMI stays under 0.2; scoped production WMI clears 0.2 on archive-complete days (shaded). Table 5: scoped production valve opens on $\approx 61\%$ OOS days ($\equiv 3/3$ ready); full-schema stress gate stays closed.

4.4 World bundle (what the LLM sees)

The consumer object is a JSON world bundle (Table 6): timing, completeness, honesty, quality (including the abstain flag), content tilts, and `evidence_ids`. Listing 1: exchange-gap day (scoped valve closed). Listing 2: 3/3-ready day where scoped WMI=1.0 opens the production valve; `full_schema_wmi` ≈ 0.093 is disclosed as ceiling audit—same day, two scopes, not a lowered q^* .

Listing 1: Exchange-gap bundle (Compiled; scoped closed).

```
{
  "date": "2026-05-11", "asset": "ADA",
  "decision_asof": "2026-05-10T23:59:00",
  "completeness": { "n_ready": 2, "n_missing": 1 },
  "world_model_index": { "band_scope": "eval_archive",
    "wmi": 0.139, "archive_complete": false,
    "should_ai_abstain": true, "thin_world": true,
    "full_schema_wmi": 0.015 },
  "macro_tilt": 1.0, "alt_tilt": -1.0
}
```



Live abstention outcomes: Sec.~Experiments (Table / four-arm figure).

Figure 1: Data-end pipeline: evidence → BandPIT → world bundle → Compiled / Ungated / Raw / Blind.




Figure 2: PIT readiness: what is visible in the archive over time.

Listing 2: 3/3-ready day: scoped valve opens.

```
{
  "date": "2026-01-16",
  "asset": "ADA",
  "completeness": {
    "n_ready": 3,
    "n_missing": 0
  },
  "world_model_index": {
    "band_scope": "eval_archive",
    "wmi": 1.0,
    "archive_complete": true,
    "should_ai_abstain": false,
    "thin_world": false,
    "full_schema_wmi": 0.093
  }
}
```




Figure 3: Dual-scope WMI on the same panel: stress ceiling vs. production contract (shaded = archive-complete).

```
"band_status": {
  "exchange": "ready",
  "macro": "ready",
  "alternative": "ready"
},
"macro_tilt": 1.0,
"alt_tilt": -1.0}
```

4.5 Four arms and frozen prompts

Actions $\in \{\text{bullish, bearish, neutral, abstain}\}$. **Compiled**: full bundle; frozen prompt *requires* abstain when should_ai_abstain is true (and when the world is judged thin). **Ungated**: same content, completeness, and numeric WMI; *no* hard flag; soft prompt asks the model to judge thinness. **Raw**: mom5 only—thin price wiring for the ladder. **Blind**: “Today is d . How should I trade a ?”

with no feed. **Dual protocol.** RQ1–RQ2 freeze *full-schema* bundles (max WMI ≈ 0.093), so Compiled abstains every OOS day—a closed stress test of the hard flag. RQ3a switches the same vendors to *scoped production* bundles: the valve opens on archive-complete days; we measure closed-day obedience, open-day action, and a parallel no-LLM tilt rule. Compiled vs. Ungated differ only in enforcement (Listings 3–4); prompts, IS/OOS cut, and seeds are frozen. On thin Compiled days the evidence-alignment proxy (EAR) is 1.0 across vendors: rationales stay bound to disclosed evidence_ids.

Listing 3: Compiled (excerpt): hard valve.

```
1. If should_ai_abstain is true OR the world is thin,
   you MUST "abstain".
2. Prefer band content over raw momentum
   (macro_tilt, alt_tilt are PIT-safe).
3. Rationale must match evidence_ids.
```

Listing 4: Ungated (excerpt): soft judgment.

```
1. Judge whether low WMI / missingness makes action
   unsafe; you MAY "abstain".
2. Prefer band content when you act.
NO hard abstain boolean; NO forced policy.
```

Why the hard valve is not a confound. Compiled abstention is partly prompt obedience, and that is intentional: the data end’s job on thin days is to *block* action, not to hope the model discovers thinness (Prop. 3.8). Ungated measures soft judgment from the same disclosed bundle; Raw measures fragment wiring; Blind measures no feed. The gating result is that the hard flag is reliable across vendors; soft refusal is not—which is why half-baked context remains dangerous without the valve.

4.6 Systems validation (engines)

Paper-lab smoke wires BandPIT/WMI/ACWMI and checks the dual gate: scoped archive opens on 3/3-ready statuses while full-schema stays closed; a handoff rule acts only when the scoped valve is open. Availability shocks O_t are queryable on the service surface.

5 Experiments

5.1 Setup

PIT panel ~ 399 days $\times 10$ liquid names (≈ 3990 asset-days; 2025-07–2026-08), chronological IS/OOS 200/200 days (cut 2026-01-16). Unless noted, live arms share the same stratified 100 OOS asset-days (seed fixed), temperature 0, and four flash/mini vendors (gpt-5.4-mini, deepseek-v4-flash, glm-5.2, gemini-3.5-flash-lite) via an OpenAI-compatible adapter. Grounding uses 50 OOS days; full-OOS abstention uses $n=2000$; no-LLM probes use ≈ 200 OOS days after return alignment. Primary metrics: RQ1 grounding F1/tilt accuracy (sufficiency chat secondary); RQ2 thin-world abstain under the *full-schema stress gate* (CE as avoided loss, not alpha); RQ3a scoped production-valve LLM obedience plus no-LLM tilt handoff; RQ3b ungated tilt Δ CE vs. momentum.

5.2 RQ1 (primary): Seeing—grounding workflow

If the data end works, models must report a *checkable* market view—not invent coverage from a price fragment.

On 50 OOS asset-days, each model lists ready/missing bands among exchange/macro/alternative, reports macro_tilt / alt_tilt

Table 7: Grounding (50 days): primary seeing result. C = Compiled; R = Raw.

Model	Ready F1		Missing F1		Tilt-sign	
	C	R	C	R	C	R
gpt-5.4-mini	0.94	0.00	0.94	0.31	0.94	0.20
deepseek-v4-flash	0.98	0.06	0.98	0.20	0.98	0.18
glm-5.2	0.98	0.09	1.00	0.11	0.98	0.20
gemini-3.5-flash-lite	0.96	0.03	0.96	0.30	0.96	0.18
Mean	0.97	0.04	0.97	0.23	0.97	0.19

Table 8: Sufficiency chat vs. field grounding (50 days). Seeing = fields; chat alone misleads.

Model	Suff. C	Suff. R	Ready F1 C	Ready F1 R
gpt-5.4-mini	0.78	0.98	0.94	0.00
deepseek-v4-flash	0.78	0.90	0.98	0.06
glm-5.2	0.80	0.96	0.98	0.09
gemini-3.5-flash-lite	0.04	0.94	0.96	0.03
Mean	0.60	0.95	0.97	0.04

signs, and judges sufficiency; (i)–(iii) score against PIT ground truth.

Table 7: Compiled mean ready/missing F1 and tilt-sign accuracy 0.97; Raw collapses to 0.04, 0.23, 0.19. Raw has no readiness or tilt fields, so it cannot recover archive state; Compiled forces use of the disclosed world.

Legibility $\neq$ sufficiency chat. Table 8: Compiled sufficiency accuracy averages only 0.60 (Gemini 0.04) despite 0.97 field grounding, while Raw often answers “enough?” correctly by chance/heuristics (0.95) while inventing coverage. Chat judgment of “is this enough to trade?” is not a substitute for a typed world; we therefore measure seeing by checkable fields (RQ1) and action safety by the hard valve (RQ2).

Case study. Listings 5–6: ADA (2026-05-11) exchange missing; Compiled recovers the checkable view; Raw cannot. A second day (2026-01-17) under the *full-schema stress* Compiled arm abstains on should_ai_abstain; Raw goes bullish from mom5 = +1 alone.

Listing 5: Compiled grounding (gpt-5.4-mini, ADA 2026-05-11).

```
ready: alternative, macro
missing: exchange
macro_tilt +; alt_tilt -; sufficiency: insufficient
"thin world... should_ai_abstain=true"
```

Listing 6: Raw grounding / action (same model family).


```
grounding: ready=[]; cannot confirm tilts
action (2026-01-17): bullish
"Positive 5-period momentum in raw feed."
```

5.3 RQ2 (safety valve): Hard refuse when the view is thin

Legibility does not imply safe action. All live LLM arms in this subsection freeze *full-schema* bundles, so every OOS day is thin

Table 9: Thin abstention on identical 100 OOS days under full-schema stress (Wilson 95% CIs). Blind = 1.0.

Model	Thin C	Thin U [CI]	Abs. R [CI]	Δ CE
gpt-5.4-mini	1.00	0.63 [.53,.72]	0.04 [.02,.10]	+0.73
deepseek-v4-flash	1.00	0.81 [.72,.87]	0.75 [.66,.82]	+1.22
glm-5.2	1.00	0.86 [.78,.91]	0.11 [.06,.19]	+1.04
gemini-3.5-flash-lite	1.00	0.43 [.34,.53]	0.07 [.03,.14]	+1.30
Mean	1.00	0.68	—	+1.07

**Figure 4: Safety valve: Blind refuses; Raw over-trades; Ungated is unstable; Compiled enforces the hard gate.****Table 10: Raw action mix (100 days): half-baked context drives directional mass.**

Model	Abs.	Bull.	Bear.	Neut.
gpt-5.4-mini	0.04	0.28	0.38	0.30
deepseek-v4-flash	0.75	0.14	0.04	0.07
glm-5.2	0.11	0.40	0.48	0.01
gemini-3.5-flash-lite	0.07	0.44	0.49	0.00

under the stress gate 0.2. Table 9 and Figure 4: on the same 100 days, Compiled abstains on every thin day (1.0); Ungated soft-judges at mean 0.68 (Gemini 0.43); Raw over-trades (Table 10); Blind abstains 1.0 with no feed—safe but useless as a data end. The costly failure mode is the middle case: half-baked Raw context makes models act and lose. Mean Δ CE (Compiled–Raw) +1.07 is avoided loss under the stress valve, not an alpha claim.

Full-OOS replication. Table 11: on all 2000 OOS asset-days (all thin under the *full-schema stress* gate 0.2), Compiled thin-abstain remains 1.000 (Wilson [.998, 1]) for every vendor. Ungated soft rates fall to a four-vendor mean ≈ 0.56 ; Raw GPT/Gemini abstain at 0.000. Scale does not rescue soft judgment. Blind refuses everywhere with no feed on the paired 100-day sample (abstain 1.0 for every vendor): the failure mode that needs a data end is not “models never abstain,” but “thin fragments make them act.”

Inside Ungated soft acts. On the mixed 100-day sample, when Ungated *does* act, act-conditional CE is sharply negative for three of four vendors (DeepSeek −3.12, GLM −4.28, Gemini −1.67) while GPT is positive on its smaller act set. Soft disclosure without a hard

Table 11: Full-OOS thin abstention ($n=2000$; all thin under full-schema stress gate 0.2).

Model	Thin C	Thin U	Abs. R
gpt-5.4-mini	1.000	0.666	0.000
deepseek-v4-flash	1.000	0.540	0.748
glm-5.2	1.000	0.679*	0.078*
gemini-3.5-flash-lite	1.000	0.359	0.000
Mean	1.000	≈ 0.56	—

*GLM coverage incomplete (Ungated $n=1402$; Raw $n=675$).

Table 12: Open slice under full-schema LLM arms (3/3 ready; $n \approx 1224$). Soft-judgment lower bound.

Model	Abs. U	CE U	Act-CE U	Abs. R	CE R
gpt-5.4-mini	0.58	+0.42	+0.32	0.00	−0.28
deepseek-v4-flash	0.42	+0.21	+0.45	0.75	+0.04
glm-5.2	0.58	+0.24	+0.11	0.09*	−0.87*
gemini-3.5-flash-lite	0.34	+0.31	+0.06	0.00	−0.25

*Raw GLM $n=398$; Ungated GLM $n=932$; others $n=1224$. Bootstrap mean CE $\approx +0.31$ ($B=5,999$ reps; one-sided $p \approx 0.21$).

flag is not a free lunch: the same readable world that RQ1 scores can still harm decisions when the valve is removed.

Open slice (still under full-schema LLM arms). Table 12: on the band-thick subset (3/3 ready $\equiv$ full-schema WMI ≥ 0.05 ; $\approx 61\%$ OOS; $n \approx 1224$) Ungated keeps compiled content without the hard boolean—mean abs. ≈ 0.48 , CE positive for every vendor (mean $\approx +0.29$; bootstrap $\approx +0.31$, CIs include 0). Full-schema Compiled still abstains at 1.0 here; Raw stays risky (GPT/Gemini abs. 0.0, CE negative). The same thick days are where the *scoped production valve* opens; live scoped LLM handoff is RQ3a.

5.4 RQ3a (handoff): Scoped production valve with live LLMs

Having established seeing (RQ1) and stress-gate refusal (RQ2), we ask whether the *production* valve is callable. The full-schema stress gate never clears 0.2, so it cannot open—by design in RQ2. Scoped production WMI opens on $\approx 61\%$ OOS days ($\equiv 3/3$ ready). We re-run the same four vendors on a stratified sample of 70 open + 30 closed OOS asset-days with scoped bundles (seed fixed; temp. 0).

Table 13: on closed days Compiled abstains 1.0 for every vendor (Ungated soft rates ≈ 0.87 – 0.97 ; Raw often keeps trading). On open days Compiled acts on a mean $\approx 89\%$ of days—the valve opens and models consume the bundle. Open-day CE does *not* beat Raw (mean Δ CE ≈ -4.2): opening the data-end valve is not an LLM alpha claim; policy skill remains downstream. Closed-day refusal avoids Raw losses (mean Δ CE $\approx +1.6$ vs. Raw on the closed slice).

A transparent no-LLM rule on the same vintaged tilts remains the clean nonempty-field handoff (Table 14): Sharpe 0.772 / CE +0.253 when the scoped valve is applied, versus momentum and buy-and-hold; full-schema stress stays fully closed.

Table 13: RQ3a live LLMs under scoped production valve (70 open + 30 closed).

Model	Abs. C_{cl}	Abs. C_{op}	CE C_{op}	CE R_{op}	ΔCE_{cl}
gpt-5.4-mini	1.00	0.03	-4.00	+0.09	+0.49
deepseek-v4-flash	1.00	0.24	-5.22	-0.21	+0.17
glm-5.2	1.00	0.11	-4.31	-0.39	+3.06
gemini-3.5-flash-lite	1.00	0.07	-4.27	-0.46	+2.64
Mean	1.00	0.11	—	—	+1.59

Abs. C_{cl}/op : Compiled abstain on closed/open days. ΔCE_{cl} = avoided loss vs. Raw on closed days.

Table 14: RQ3a no-LLM handoff. Scoped production valve vs. full-schema stress.


Policy (reads bundle fields)	Sharpe	CE [†]	Abstain
Full-schema stress 0.2	0.000	0.000	1.00
Scoped production 0.2 + tilts	0.772	+0.253	0.39
Momentum (mom5)	0.101	-0.075	0.00
Buy-and-hold	-1.399	-0.987	0.00

[†]CE = ann. mean - $\frac{1}{2}$ ann. variance (valve-aware path).

Table 15: RQ3b ungated content probe (no valve). Bootstrap ΔCE 0.334.

Policy	Sharpe	CE [‡]
Buy-and-hold	-1.404	-1.163
Momentum	0.096	-0.206
World-bundle tilts	0.764	+0.130
Tilts - Momentum	—	0.334 ($p=0.034$)

[‡]Content-harness CE (always-on; not comparable 1:1 to [†]).

**Figure 5: Secondary content probe: compiled tilts beat momentum; ΔCE positive for every asset.**

5.5 RQ3b (content): Nonempty fields

Table 15 and Figure 5: ungated always-on tilts vs. momentum on the same vintaged fields ($\Delta CE=0.334$, $p=0.034$; LOBO content share 1.0). This probe shuts the gating channel (Prop. 3.9); it is not the production valve.


Table 16: RQ3b content probe by archive density (OOS; ungated).

Slice	Share	Mech CE	ΔCE vs mom
All OOS	1.00	0.130	0.334*
Open / thick	0.61	0.061	0.315
Thin (exchange gap)	0.39	1.048 [§]	0.129

*Block-bootstrap ΔCE ; raw mean gap 0.336. [§]Always-on selection artifact.

Table 17: LOBO decomposition (ungated content probe). Content share 1.0.

Drop	ΔCE content	p	Content share
macro	-0.334	0.034	1.0
alternative	-0.334	0.034	1.0
exchange	0.000	1.000	—

**Figure 6: LOBO: deleting macro/alternative removes the edge via content, not the refuse flag.**

Density split and LOBO channel. Table 16: open/thick $\Delta CE=0.315$; on exchange-gap days the always-on probe still trades (high slice CE is a selection artifact, not a valve claim; $\Delta CE=0.129$). Table 17 and Figure 6 (Prop. 3.9): deleting macro or alternative collapses the gap through the content channel (share 1.0; gating residual 0).

6 Discussion

What carries the paper. (i) As an AI-finance data substrate, the compiled world is legible: grounding ≈ 0.97 , with sufficiency that a poor substitute (RQ1 / Props. 3.6–3.7). (ii) Under the full-schema stress gate, soft text fails at scale; Blind refuses with no feed; Raw acts on fragments and loses; the hard flag works (RQ2 / Prop. 3.8). (iii) Scoped production WMI opens on band-thick days; live LLMs abstain 1.0 when closed and act when open, while open-day CE does not beat Raw—observation usability $\neq$ policy alpha (RQ3a / Def. 3.4+Prop. 3.8); a no-LLM tilt rule and LOBO content share 1.0 keep fields nonempty (RQ3b / Prop. 3.9). Relative to FinMem /FinAgent and ICAIF agent stacks [1, 3, 15, 16, 20, 21], we supply—and evaluate—the observation stage those systems assume: data for agents, and a yardstick for whether agents understood the market.

Architecture. Semantics (Section 3) define the interfaces; Runtime (Section 4) implements Π_t and the service handoff; Experiments (Section 5) test seeing, stress gating, then scoped handoff and content in that order. Trading PnL is intentionally secondary: it confounds observation quality with policy skill.

Design. World bundle not feature dump; previous-close PIT with explicit freshness; four arms separate seeing+gating, seeing without gating, thin wiring, and no feed. Dual WMI scope separates stress audit from production readiness; health/bundle/handoff reads make both callable. Evaluative objects stay world-centric: grounding and evidence-alignment (EAR) for understanding; thin-abstain for honesty; open-day usability without claiming a predictive Pareto win.

Reproducibility. Paper-lab smoke checks BandPIT/WMI/ACWMI and scoped-vs-full gates; PIT rebuild and the frozen four-arm harness regenerate the panel and live transcripts used here. Prompts, IS/OOS cut, seeds, and vendor IDs are fixed; credentials stay in environment variables. An anonymized harness ships on acceptance.

Estimand and scope. This paper measures data-end legibility, stress-gate reliability, and whether a scoped production valve is obeyed and callable. RQ1–RQ2 isolate seeing and hard refusal under the full-schema stress protocol; RQ3a switches to scoped production bundles for live LLM obedience and a parallel no-LLM tilt rule. Open-day LLM CE is reported as a usability check, not an alpha claim—policy skill sits downstream of the data end.

7 Conclusion

AI-in-finance LLM/agent stacks need a reliable data end before decision or execution. We contribute that observation layer—data world-model infrastructure for market analysis: compile asynchronous evidence into a consumable PIT world bundle, disclose structural missingness, and hard-refuse when the view is too thin. Grounding recovers ready/missing/tilt at ≈ 0.97 under Compiled vs. Raw collapse; sufficiency chat alone fails; the full-schema stress gate abstains 1.0 where soft judgment does not; under the scoped production valve live LLMs abstain 1.0 on closed days and act when open, while a no-LLM tilt rule stays nonempty (Sharpe 0.77). Move from model-first to data-end-first: compile so models can see; refuse when they must not act.

References

- [1] Chanyeol Choi, Jihoon Kwon, Alejandro Lopez-Lira, Chaewoon Kim, Minjae Kim, Juneha Hwang, Jaeseon Ha, Hojun Choi, Suyeol Yun, Yongjin Kim, and Yongjae Lee. 2025. FinAgentBench: A Benchmark Dataset for Agentic Retrieval in Financial Question Answering. In *Proceedings of the 6th ACM International Conference on AI in Finance*. ACM, 632–637. doi:10.1145/3768292.3770362
- [2] Dean Croushore and Tom Stark. 2001. A Real-Time Data Set for Macroeconomists. *Journal of Econometrics* 105, 1 (2001), 111–130.
- [3] Yitong Duan, Chuheng Zhang, and Jian Li. 2025. FactorMAD: A Multi-Agent Debate Framework Based on Large Language Models for Interpretable Stock Alpha Factor Mining. In *Proceedings of the 6th ACM International Conference on AI in Finance*. ACM, 605–613.
- [4] Ran El-Yaniv and Yair Wiener. 2010. On the Foundations of Noise-Free Selective Classification. *Journal of Machine Learning Research* 11 (2010), 1605–1641.
- [5] Yonatan Geifman and Ran El-Yaniv. 2017. Selective Classification for Deep Neural Networks. In *Advances in Neural Information Processing Systems*, Vol. 30. Curran Associates.
- [6] David Ha and Jürgen Schmidhuber. 2018. Recurrent World Models Facilitate Policy Evolution. In *Advances in Neural Information Processing Systems*, Vol. 31. Curran Associates.
- [7] Danijar Hafner, Timothy Lillicrap, Jimmy Ba, and Mohammad Norouzi. 2020. Dream to Control: Learning Behaviors by Latent Imagination. In *International Conference on Learning Representations*. OpenReview.net.
- [8] Campbell R. Harvey, Yan Liu, and Heqing Zhu. 2016. ... and the Cross-Section of Expected Returns. *Review of Financial Studies* 29, 1 (2016), 5–68.
- [9] Yann LeCun. 2022. A Path Towards Autonomous Machine Intelligence. *OpenReview preprint* (2022).
- [10] Yang Li, Yangyang Yu, Haohang Li, Zhiyang Chen, and Khaldoun Khashanah. 2023. TradingGPT: Multi-Agent System with Layered Memory and Distinct Characters for Enhanced Financial Trading Performance. *arXiv preprint arXiv:2309.03736* (2023).
- [11] Yukun Liu, Aleh Tsyvinski, and Xi Wu. 2022. Common Risk Factors in Cryptocurrency Markets. *Journal of Finance* 77, 2 (2022), 1133–1177.
- [12] Alejandro Lopez-Lira and Yuehua Tang. 2023. Can ChatGPT Forecast Stock Price Movements? Return Predictability and Large Language Models. *arXiv preprint arXiv:2304.07619* (2023).
- [13] Igor Makarov and Antoinette Schoar. 2020. Trading and Arbitrage in Cryptocurrency Markets. *Journal of Financial Economics* 135, 2 (2020), 293–319.
- [14] Dimitris N. Politis and Joseph P. Romano. 1994. The Stationary Bootstrap. *J. Amer. Statist. Assoc.* 89, 428 (1994), 1303–1313.
- [15] Yiqing Shen, Jingshu Zhang, Feng Chen, Kaiyuan Yan, and Hongguang Li. 2025. FinSearch: A Temporal-Aware Search Agent Framework for Real-Time Financial Information Retrieval with Large Language Models. In *Proceedings of the 6th ACM International Conference on AI in Finance*. ACM. doi:10.1145/3768292.3770382
- [16] Rui Sun, Zuo Bai, Wentao Zhang, Yuxiang Zhang, Li Zhao, Shan Sun, and Zhengwen Qiu. 2025. FinResearchBench: A Logic Tree based Agent-as-a-Judge Evaluation Framework for Financial Research Agents. In *Proceedings of the 6th ACM International Conference on AI in Finance*. ACM, 656–664. doi:10.1145/3768292.3770364
- [17] Halbert White. 2000. A Reality Check for Data Snooping. *Econometrica* 68, 5 (2000), 1097–1126.
- [18] Hongyang Yang, Xiao-Yang Liu, and Christina Dan Wang. 2023. FinGPT: Open-Source Financial Large Language Models. In *FinLLM Symposium at IJCAI*.
- [19] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2023. ReAct: Synergizing Reasoning and Acting in Language Models. In *International Conference on Learning Representations*. OpenReview.net.
- [20] Yangyang Yu, Haohang Li, Zhiyang Chen, Yuechen Jiang, Yang Li, Denghui Zhang, Rong Liu, Jordan W. Suchow, and Khaldoun Khashanah. 2024. FinMem: A Performance-Enhanced LLM Trading Agent with Layered Memory and Character Design. *Proceedings of the AAAI Symposium Series* 3, 1 (2024), 595–597.
- [21] Wentao Zhang, Lingxuan Zhao, Hao Xia, Shuo Sun, Jiaze Sun, Molei Qin, Xinyi Li, Yuqing Zhao, Yilei Zhao, Xinyu Cai, et al. 2024. A Multimodal Foundation Agent for Financial Trading: Tool-Augmented, Diversified, and Generalist. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*. ACM.